

IT5002

Computer Systems and Applications

Caches

colintan@nus.edu.sg

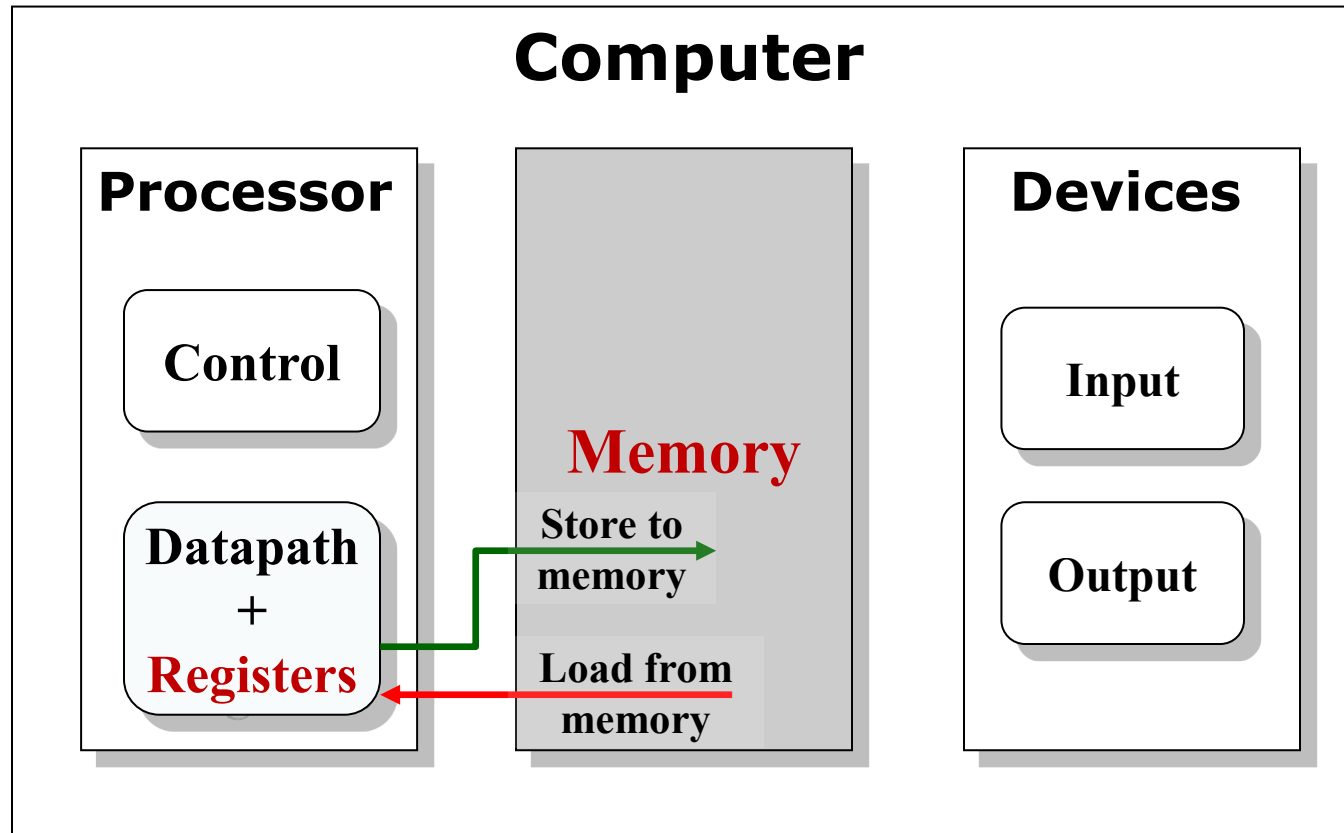


Q & A

- **DO NOT** use the Zoom chat for questions. It doesn't appear in the video recordings.
- Please ask questions at <https://sets.netlify.app/module/61597486a7805d9fb1b4accd>

← OR scan this QR code (may be obscured on some slides)

1. Data Transfer: The Big Picture



Registers are in the datapath of the processor. If operands are in memory we have to **load** them to processor (registers), operate on them, and **store** them back to memory.


1. Memory Technology: 1950s



1948: Maurice Wilkes examining EDSAC's delay line memory tubes 16-tubes each storing 32 17-bit words



Maurice Wilkes: 2005



1952: IBM 2361 16KB magnetic core memory




Figure 10. IBM 2361 Core Storage

1. Memory Technology Today: DRAM

■ **DDR SDRAM**

- **Double Data Rate**
 - Synchronous **D**ynamic **R**AM
- The dominant memory technology in PC market
- Delivers memory on the positive and negative edge of a clock (double rate)
- Generations:
 - **DDR** ($\text{MemClkFreq} \times 2(\text{double rate}) \times 8 \text{ words}$)
 - **DDR2** ($\text{MemClkFreq} \times 2(\text{multiplier}) \times 2 \times 8 \text{ words}$)
 - **DDR3** ($\text{MemClkFreq} \times 4(\text{multiplier}) \times 2 \times 8 \text{ words}$)
 - **DDR4** (Lower power consumption, higher bandwidth)




1. DRAM Capacity Growth


Growth of Capacity per DRAM Chip

❖ DRAM capacity quadrupled almost every 3 years

✧ 60% increase per year, for 20 years



DRAM Chip Capacity




- Unprecedented growth in density, but we still have a problem

1. Processor-DRAM Performance Gap


Memory Wall:

1GHz Processor → 1 ns per clock cycle

50ns for DRAM access → 50 processor clock cycles per memory access!



1. Faster Memory Technology: SRAM



SRAM

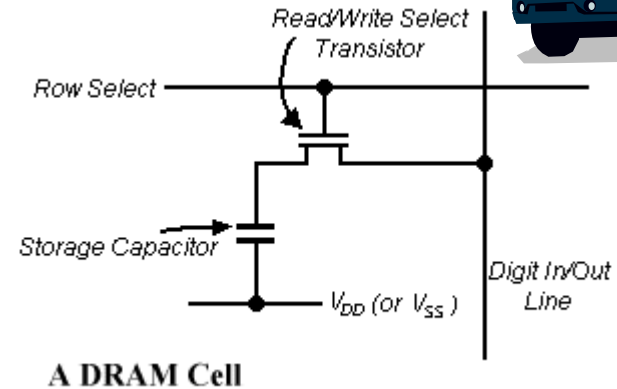
6 transistors per memory cell

→ **Low density**

Fast access latency of 0.5 – 5 ns

More costly

Uses flip-flops



DRAM

1 transistor per memory cell

→ **High density**


Slow access latency of 50-70ns

Less costly

Used in main memory

1. Slow Memory Technology: Magnetic Disk

Drive Physical and Logical Organization




Typical high-end hard disk:

Average Latency: 4 - 10 ms

Capacity: 500-2000GB

1. Quality vs Quantity



	Capacity	Latency	Cost/GB
Register	100s Bytes	20 ps	\$\$\$\$
SRAM	100s KB	0.5-5 ns	\$\$\$
DRAM	100s MB	50-70 ns	\$
Hard Disk	100s GB	5-20 ms	Cents
Ideal	1 GB	1 ns	Cheap

1. Best of Both Worlds


- What we want:
 - A **BIG** and **FAST** memory
 - Memory system should perform like 1GB of SRAM (1ns access time) but cost like 1GB of slow memory

Key concept:

Use a **hierarchy** of memory technologies:

- ❖ Small but fast memory near CPU
- ❖ Large but slow memory farther away from CPU

1. Memory Hierarchy



2. Cache: The Library Analogy



Imagine you are forced to put back a book to its bookshelf before taking another book.....

2. Solution: Book on the Desk!



What if you are allowed to take the books that are **likely to be needed soon** with you and place them nearby on the desk?

2. Cache: The Basic Idea

- **Keep the frequently and recently used data in smaller but faster memory**
- **Refer to bigger and slower memory:**
 - **Only when you cannot find data/instruction in the faster memory**
- **Why does it work?**

Principle of Locality

Program accesses only a small portion of the memory address space within a small time interval

2.1 Cache: Types of Locality

■ Temporal locality


- If an item is referenced, it will tend to be referenced again soon

■ Spatial locality

- If an item is referenced, nearby items will tend to be referenced soon

■ Different locality for


- Instructions
- Data




2.1 Working Set: Definition

- Set of locations accessed during Δt
- Different phases of execution may use different working sets

Our aim is to **capture the working set and keep it in the memory closest to CPU**




2.2 Two Aspects of Memory Access



- How to make SLOW main memory appear faster?
 - **Cache** – a small but fast SRAM near CPU
 - **Hardware managed:** Transparent to programmer
- How to make SMALL main memory appear bigger than it is?
 - **Virtual memory**
 - **OS managed:** Transparent to programmer
 - Not in the scope of this module (covered in CS2106)

2.2 Memory Access Time: Terminology



- **Hit:** Data is in cache (e.g., X)
 - **Hit rate:** Fraction of memory accesses that hit
 - **Hit time:** Time to access cache
- **Miss:** Data is not in cache (e.g., Y)
 - **Miss rate** = $1 - \text{Hit rate}$
 - **Miss penalty:** Time to replace cache block + hit time
- **Hit time** < **Miss penalty**

2.2 Memory Access Time: Formula

Average Access Time

$$= \text{Hit rate} \times \text{Hit Time} + (1 - \text{Hit rate}) \times \text{Miss penalty}$$

Example:

- Suppose our on-chip SRAM (cache) has **0.8 ns** access time, but the fastest DRAM (main memory) we can get has an access time of **10ns**. How high a hit rate do we need to sustain an average access time of **1ns**?

Let h be the desired hit rate.

$$1 = 0.8h + (1 - h) \times (10 + 0.8)$$

$$= 0.8h + 10.8 - 10.8h$$

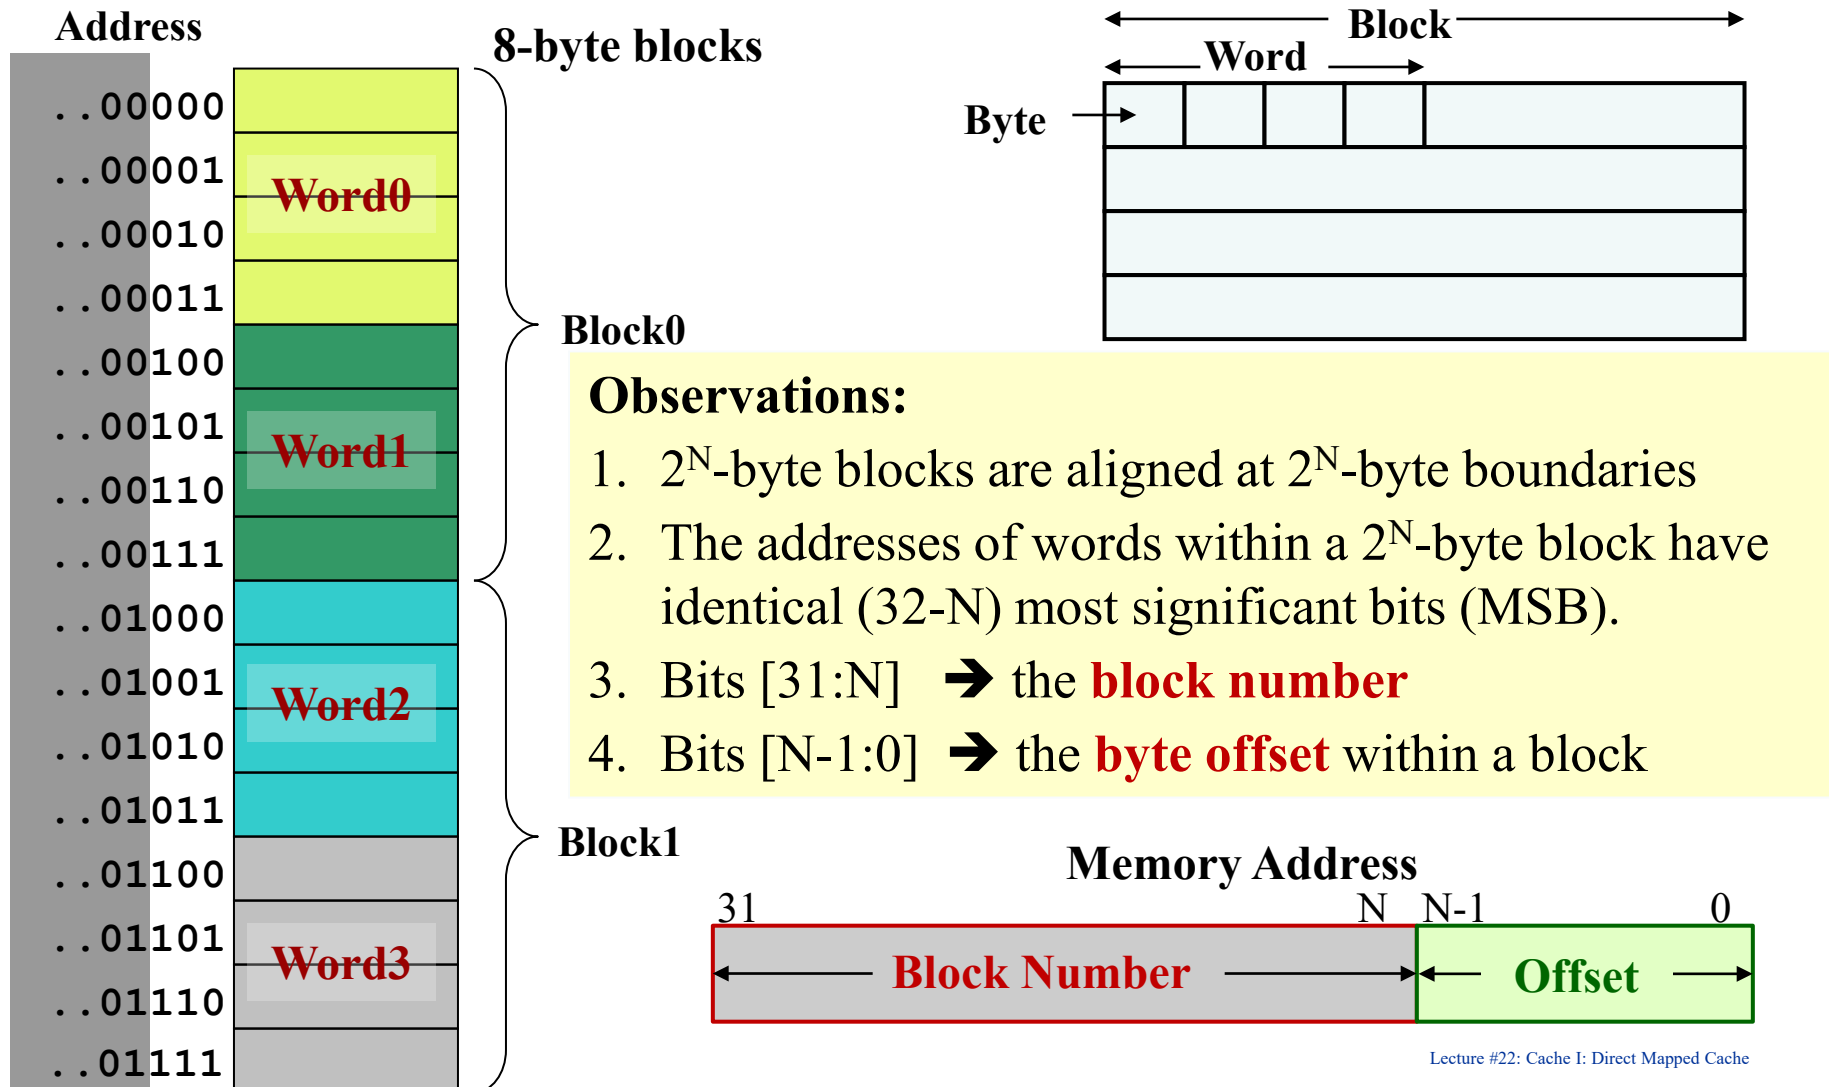
$$10h = 9.8 \rightarrow h = 0.98$$

Hence we need a hit rate of **98%**.

3. Memory to Cache Mapping (1/2)

- **Cache Block/Line:**
 - Unit of transfer between memory and cache
- Block size is typically one or more words
 - e.g.: 16-byte block \cong 4-word block
 - 32-byte block \cong 8-word block
- Why is the block size bigger than word size?

3. Memory to Cache Mapping (2/2)

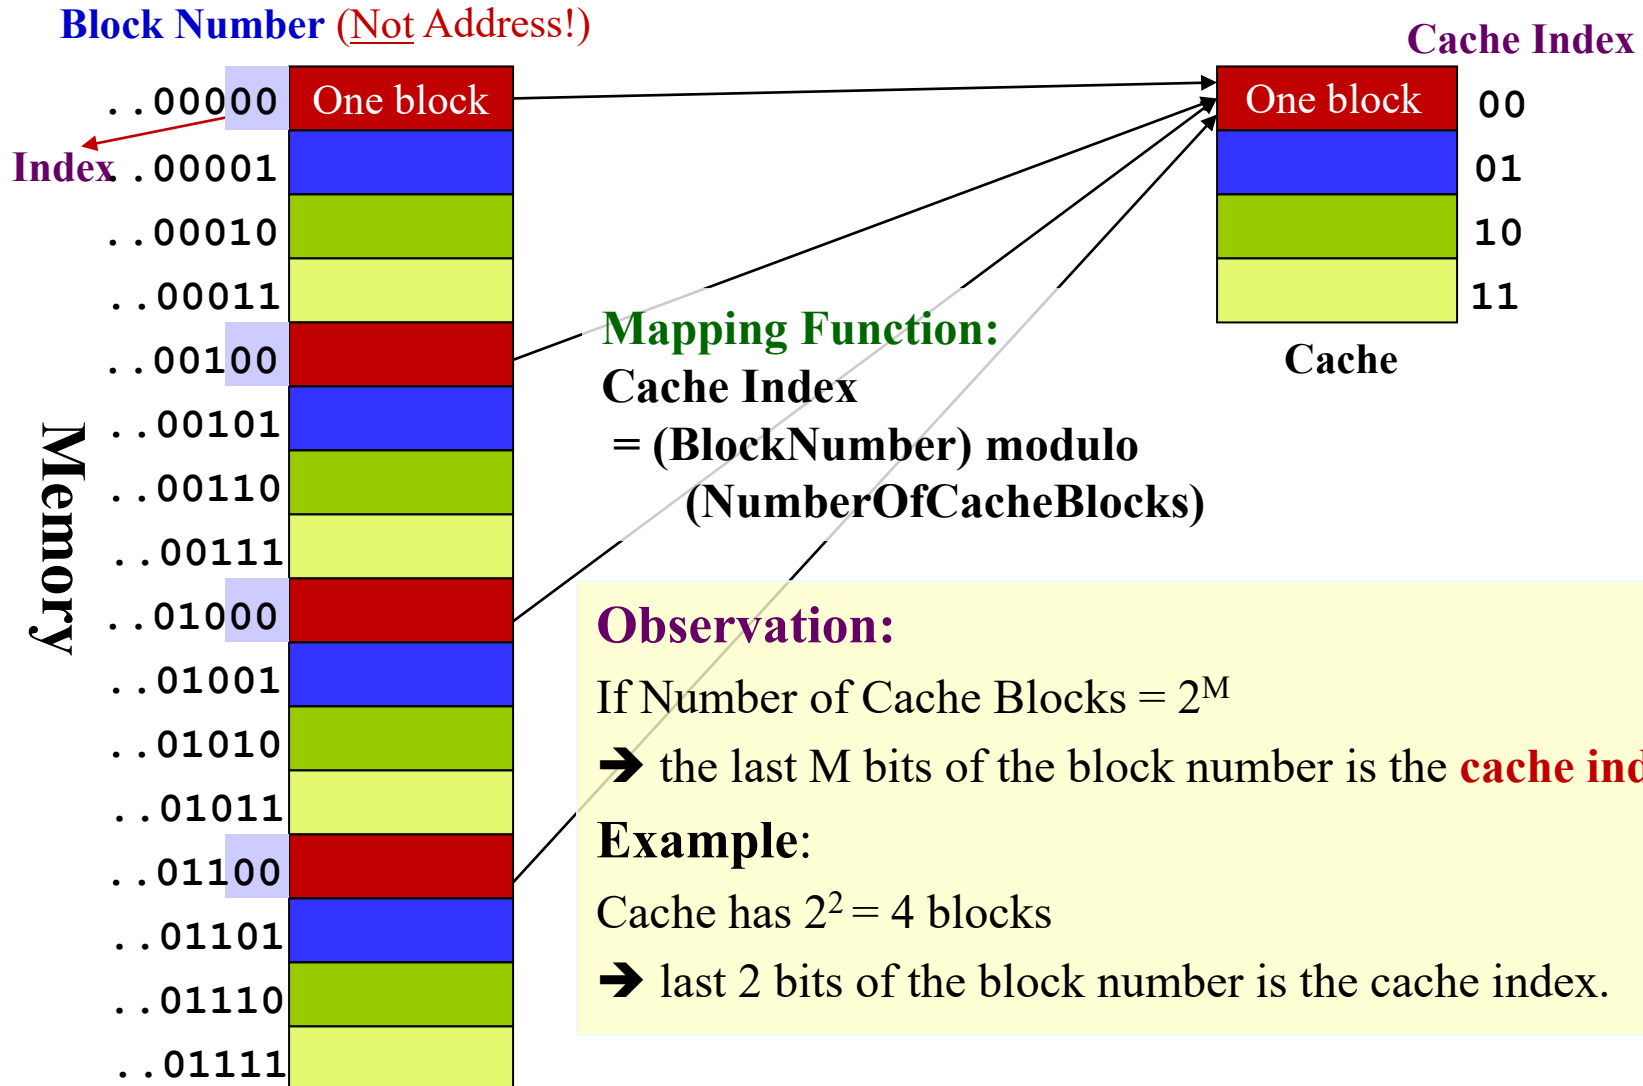


4. Direct Mapping Analogy

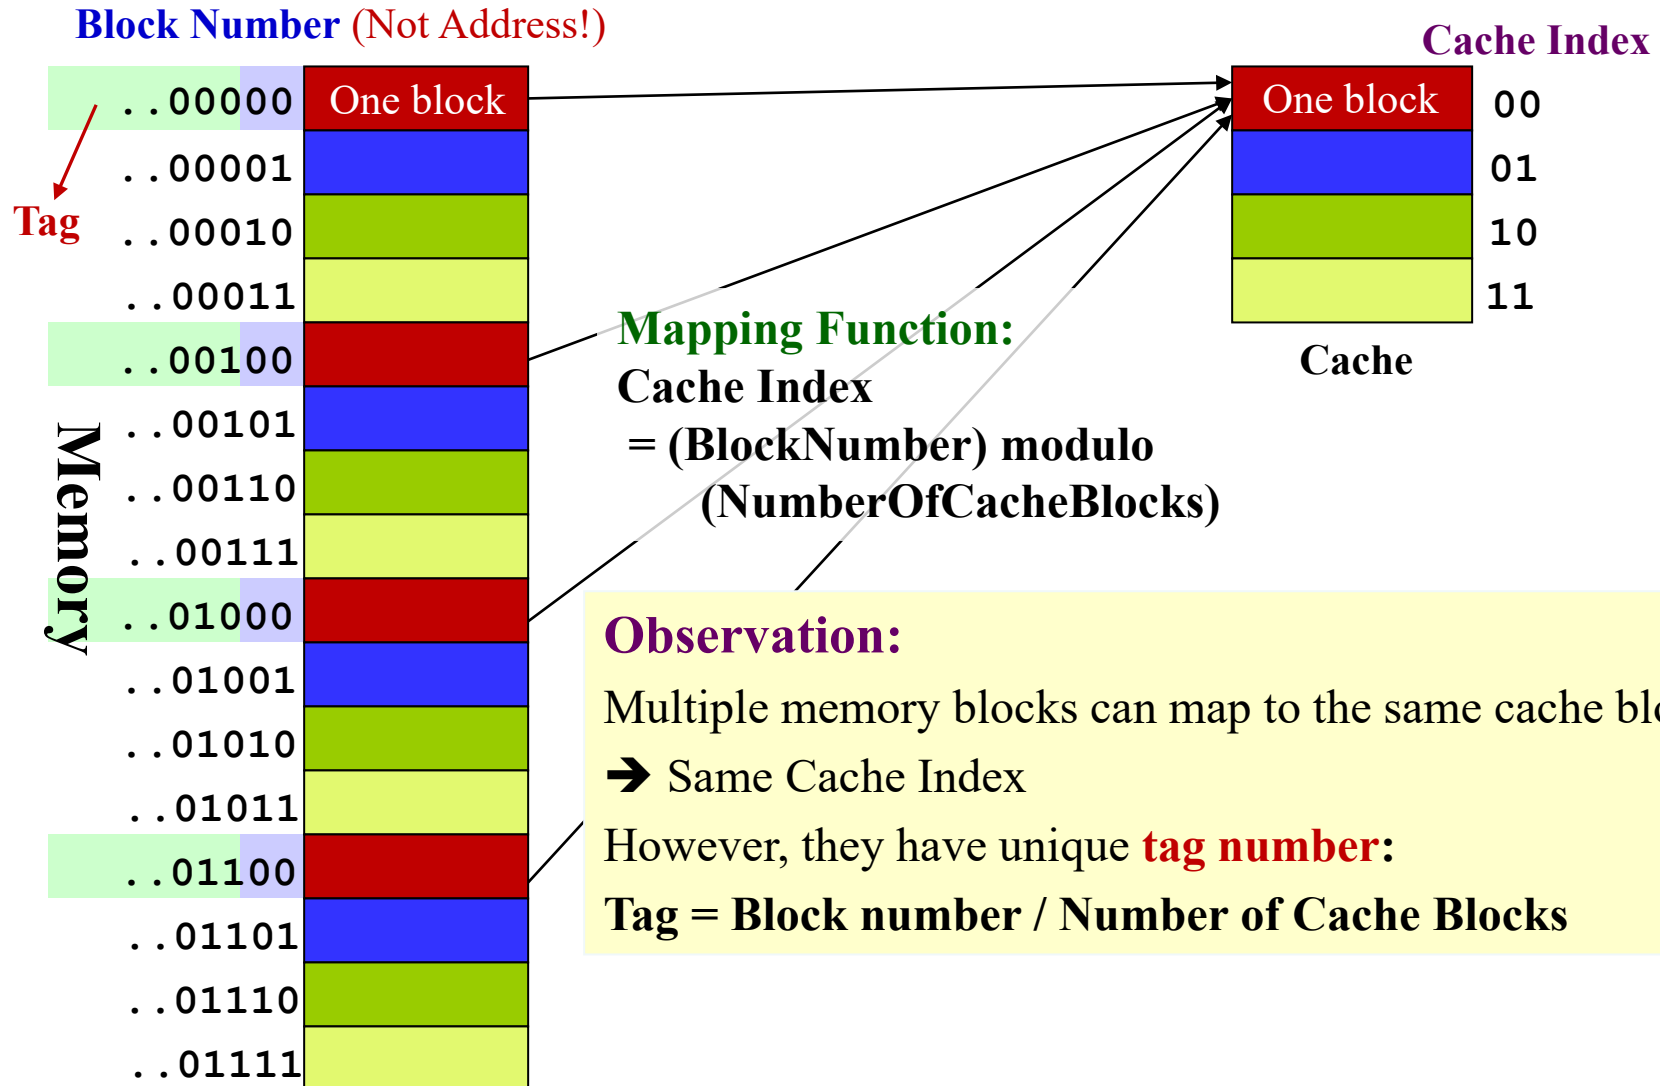


Imagine there are 26 “locations” on the desk to store books. A book’s location is determined by the first letter of its title.
➔ Each book has exactly one location.


4. Direct Mapped Cache: Cache Index




4. Direct Mapped Cache: Cache Tag



4. Direct Mapped Cache: Mapping



Cache Block size = 2^N bytes



Cache Block size = 2^N bytes

Number of cache blocks = 2^M

Offset = N bits

Index = M bits

Tag = $32 - (N + M)$ bits

4. Direct Mapped Cache: Cache Structure




Cache	Valid	Tag	Data	Index
				00
				01
				10
				11

Along with a data block (line), cache also contains the following administrative information (overheads):

1. **Tag** of the memory block
2. **Valid bit** indicating whether the cache line contains valid data

When is there a cache hit?
 (Valid[index] = TRUE) AND
 (Tag[index] = Tag[memory address])


4. Cache Mapping: Example



Offset, **N** = 4 bits

Block Number = $32 - 4 = 28$ bits

Check: Number of Blocks = 2^{28}



Number of Cache Blocks

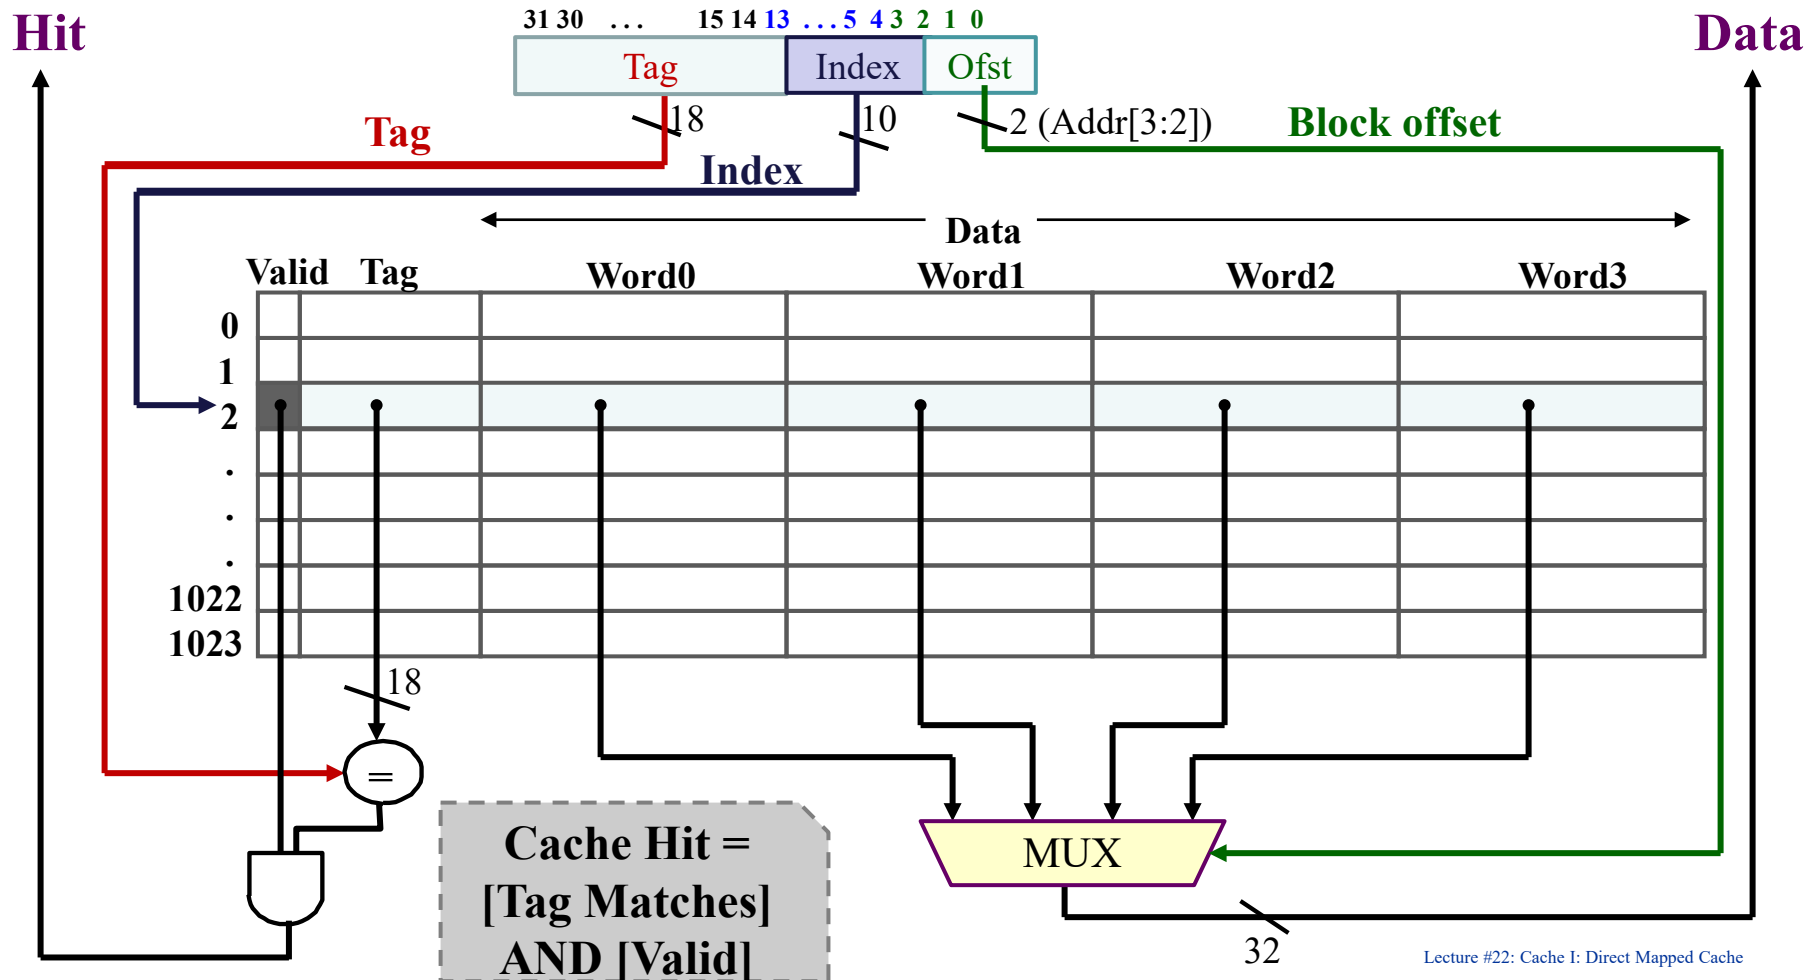
$$= 16\text{KB} / 16\text{bytes} = 1024 = 2^{10}$$

Cache Index, **M** = 10bits

Cache Tag = $32 - 10 - 4 = 18$ bits

4. Cache Circuitry: Example

16-KB cache:
4-word (16-byte) blocks



5. Reading Data: Setup

- **Given a direct mapped 16KB cache:**
 - 16-byte blocks x 1024 cache blocks
- **Trace the following memory accesses:**

Tag														Index				Offset		
31														14	13		4	3	0	
000000000000000000000000														0000000001				0100		
000000000000000000000000														0000000001				1100		
000000000000000000000000														0000000011				0100		
000000000000000000000010														0000000001				1000		
000000000000000000000000														0000000001				0000		

5. Reading Data: Initial State


- Initially cache is empty
 → All *valid* bits are zeroes (false)

		← Data →																	
		Valid	Tag	Word0 Bytes 0-3				Word1 Bytes 4-7				Word2 Bytes 8-11				Word3 Bytes 12-15			
0		0																	
1		0																	
2		0																	
3		0																	
4		0																	
5		0																	
						
1022		0																	
1023		0																	

5. Reading Data: Load #1-1

Tag**Index****Offset**

▪ Load from



Step 1. Check Cache Block at index 1

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #1-2

Tag**Index****Offset**

- Load from

00000000000000000000

0000000001

0100

Step 2. Data in block 1 is invalid [Cold/Compulsory Miss]

← Data →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #1-3

- Load from

Tag	Index	Offset
00000000000000000000	0000000001	0100

Step 3. Load 16 bytes from memory; Set Tag and Valid bit

		← Data →					
		Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0		0					
1		1	0	A	B	C	D
2		0					
3		0					
4		0					
5		0					
...							
1022		0					
1023		0					

5. Reading Data: Load #1-4

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000001

0100



Step 4. Return Word1 (byte offset = 4) to Register

← **Data** →

Index	Valid	Tag	Word0	Word1	Word2	Word3
			Bytes 0-3	Bytes 4-7	Bytes 8-11	Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #2-1

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000001

1100

Step 1. Check Cache Block at index 1

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #2-2

- **Load from**
- | Tag | Index | Offset |
|----------------------|------------|--------|
| 00000000000000000000 | 0000000001 | 1100 |
- Step 2. [Cache Block is Valid] AND [Tags match] → **Cache hit!**

← Data →						
Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #2-3

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000001

1100

Step 3. Return Word3 (byte offset = 12) to Register [Spatial Locality]

← **Data** →

			Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
Index	Valid	Tag				
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #3-1

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000011

0100

Step 1. Check Cache Block at index 3

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #3-2

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000011

0100

Step 2. Data in block 3 is invalid [Cold/Compulsory Miss]

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	0					
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #3-3

- Load from

Tag	Index	Offset
00000000000000000000	0000000011	0100

Step 3. Load 16 bytes from memory; Set Tag and Valid bit

		Data			
		Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
Index	Valid	Tag			
0	0				
1	1	0	A	B	C
2	0				
3	1	0	I	J	K
4	0				
5	0				
...					
1022	0				
1023	0				

5. Reading Data: Load #3-4

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000011

0100



Step 4. Return Word1 (byte offset = 4) to Register

← **Data** →

			Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
Index	Valid	Tag				
0	0					
1	1	0	A	B	C	D
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #4-1

Tag**Index****Offset**

▪ Load from

0000000000000000010

0000000001

1000

Step 1. Check Cache Block at index 1

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #4-2

- Load from

Tag	Index	Offset
00000000000000000010	0000000001	1000
- Step 2. Cache block is Valid but Tags mismatch [Cold miss]

		← Data →				
Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #4-3

Tag**Index****Offset**

▪ Load from

0000000000000000010

0000000001

1000

Step 3. Replace block 1 with new data; Set Tag

				Data			
				Word0	Word1	Word2	Word3
				Bytes 0-3	Bytes 4-7	Bytes 8-11	Bytes 12-15
Index	Valid	Tag					
0	0						
1	1	2		E	F	G	H
2	0						
3	1	0		I	J	K	L
4	0						
5	0						
...							
1022	0						
1023	0						

5. Reading Data: Load #4-4

Tag**Index****Offset**

▪ Load from

0000000000000000010

0000000001

1000



Step 4. Return Word2 (byte offset = 8) to Register

← Data →

			Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
Index	Valid	Tag				
0	0					
1	1	2	E	F	G	H
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #5-1

Tag**Index****Offset**

▪ Load from



Step 1. Check Cache Block at index 1

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	2	E	F	G	H
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
...						
1022	0					
1023	0					

5. Reading Data: Load #5-2

Tag**Index****Offset**

▪ Load from



Step 2. Cache block is Valid but Tags mismatch [**Cold miss**]

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	2	E	F	G	H
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #5-3

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000001

0000

Step 3. Replace block 1 with new data; Set Tag

← **Data** →

Index	Valid	Tag	Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
0	0					
1	1	0	A	B	C	D
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Load #5-4

Tag**Index****Offset**

▪ Load from

00000000000000000000

0000000001

0000

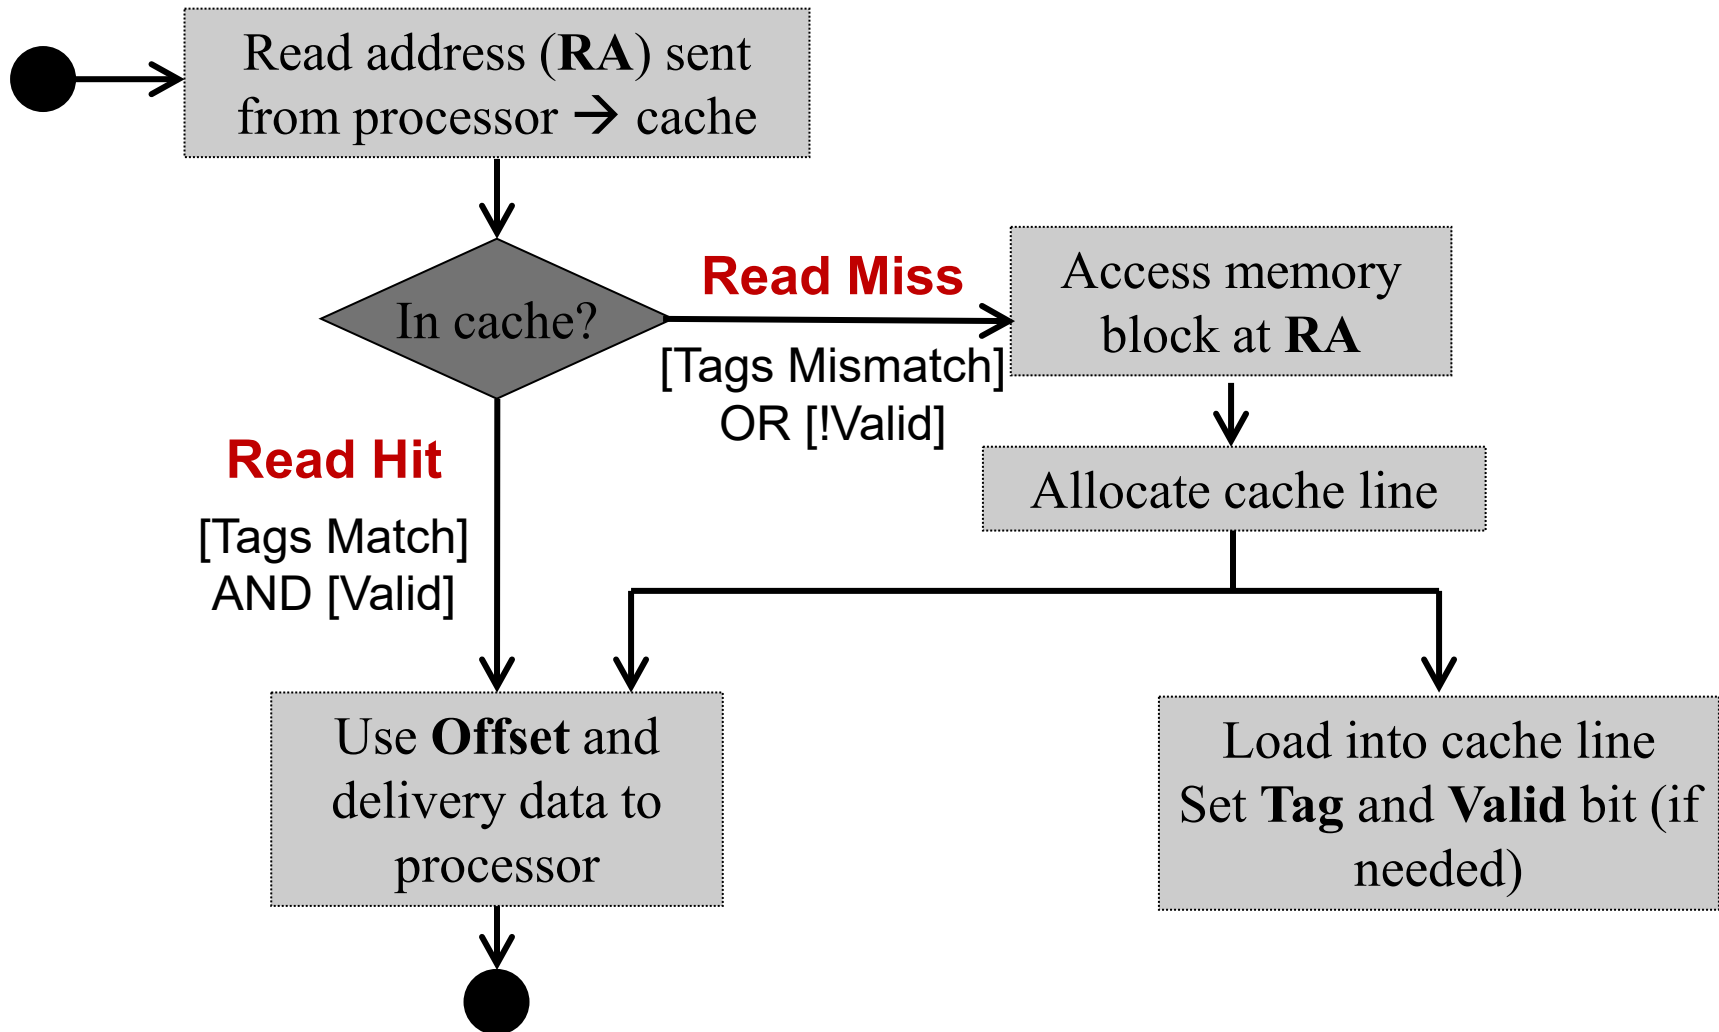


Step 4. Return Word0 (byte offset = 0) to Register

← **Data** →

			Word0 Bytes 0-3	Word1 Bytes 4-7	Word2 Bytes 8-11	Word3 Bytes 12-15
Index	Valid	Tag				
0	0					
1	1	0	A	B	C	D
2	0					
3	1	0	I	J	K	L
4	0					
5	0					
... ..						
1022	0					
1023	0					

5. Reading Data: Summary



6. Types of Cache Misses

■ **Compulsory misses**

- On the first access to a block; the block must be brought into the cache
- Also called **cold start misses** or **first reference misses**

■ **Conflict misses**

- Occur in the case of direct mapped cache or set associative cache, when several blocks are mapped to the same block/set
- Also called **collision misses** or **interference misses**


■ **Capacity misses**

- Occur when blocks are discarded from cache as cache cannot contain all blocks needed

7. Changing Cache Content: **Write Policy**

- Cache and main memory are inconsistent
 - Modified data only in cache, not in memory!
- **Solution 1: Write-through** cache
 - Write data both to cache and to main memory
- **Solution 2: Write-back** cache
 - Only write to cache
 - Write to main memory only when cache block is replaced (evicted)

7. Write-Through Cache



- **Problem:**
 - Write will operate at the speed of main memory!
- **Solution:**
 - Put a write buffer between cache and main memory
 - Processor: writes data to cache + write buffer
 - Memory controller: write contents of the buffer to memory

7. Write-Back Cache

- **Problem:**

- Quite wasteful if we write back every evicted cache blocks

- **Solution:**

- Add an additional bit (**Dirty bit**) to each cache block
- Write operation will change dirty bit to 1
 - Only cache block is updated, no write to memory
- When a cache block is replaced:
 - Only write back to memory if dirty bit is 1

7. Handling Cache Misses

- On a **Read Miss**:

- Data loaded into cache and then load from there to register

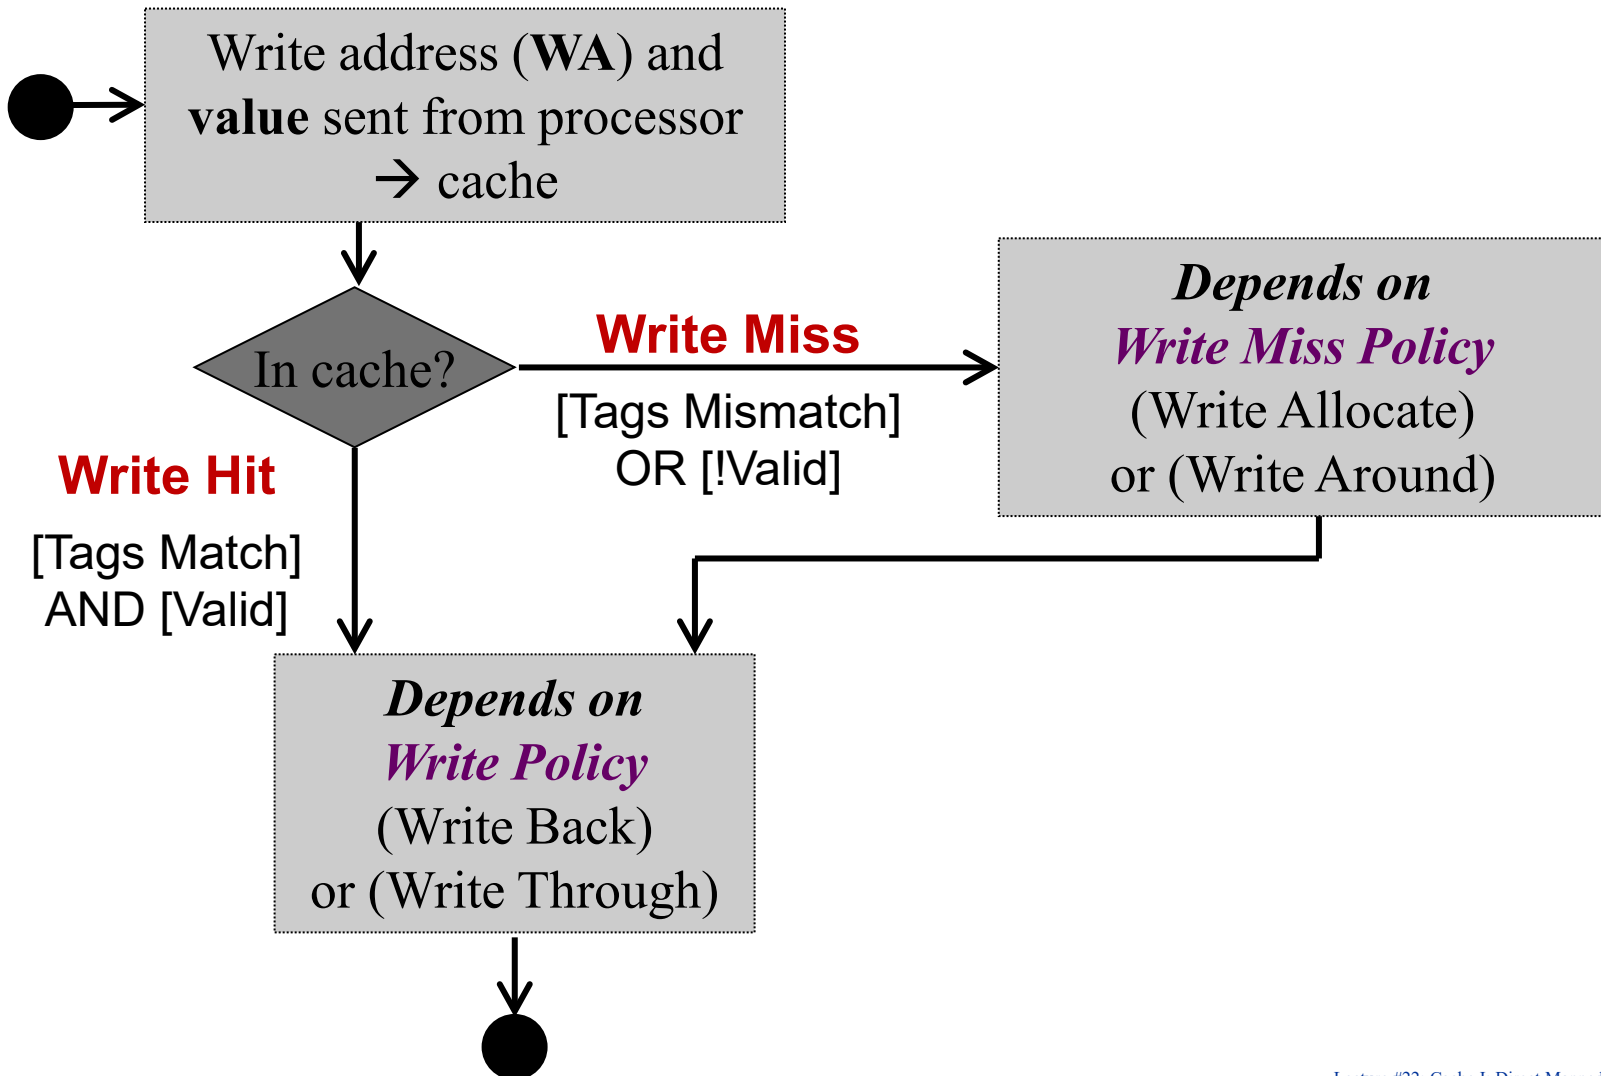
- **Write Miss** option 1: **Write allocate**

- Load the complete block into cache
- Change only the required word in cache
- Write to main memory depends on write policy

- **Write Miss** option 2: **Write around**

- Do not load the block to cache
- Write directly to **main memory only**

7. Writing Data: Summary



8. Set Associative (SA) Cache

■ Compulsory misses

- On the first access to a block, it is brought into the cache
- Also called cold start misses

Solution:
Set Associative Cache

■ Conflict misses

- Occur in the case of direct mapped cache or set associative cache, when several blocks are mapped to the same block/set
- Also called **collision misses** or **interference misses**

■ Capacity misses

- Occur when blocks are discarded from cache as cache cannot contain all blocks needed

8. Set Associative Cache: Analogy



Many book titles start with “T”

→ Too many conflicts!

Hmm... how about we give more slots
per letter, 2 books start with “A”, 2
books start with “B”, etc?

8. Set Associative (SA) Cache

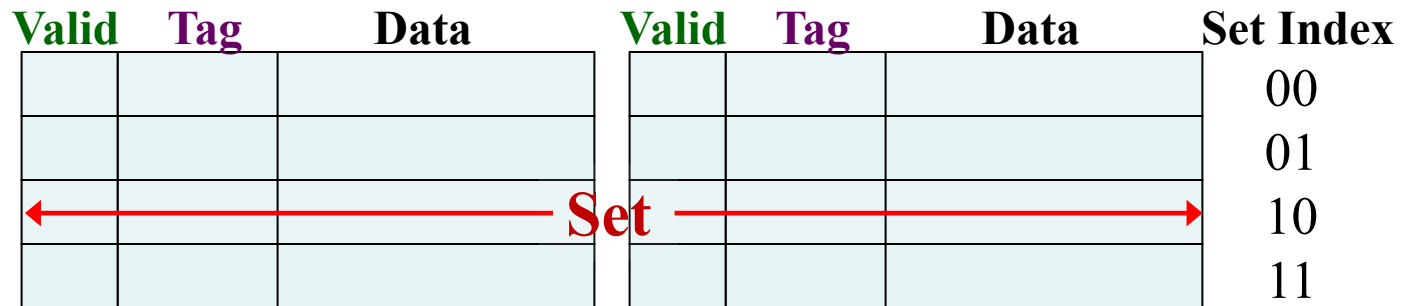
- **N-way Set Associative Cache**

- A memory block can be placed in a fixed number (**N**) of locations in the cache, where **N** > 1

- **Key Idea:**

- Cache consists of a number of sets:
 - **Each set contains N cache blocks**
- Each memory block maps to a unique cache set
- Within the set, a memory block can be placed in **any** of the **N** cache blocks in the set


8. Set Associative Cache: Structure



2-way Set Associative Cache

- **An example of 2-way set associative cache**
 - Each set has two cache blocks
 - **A memory block maps to a unique set**
 - In the set, the memory block can be placed in **either of the cache blocks**
- ➔ Need to search both to look for the memory block


8. Set Associative Cache: Mapping



Cache Block size = 2^N bytes

Cache Set Index

= (BlockNumber) modulo (NumberOfCacheSets)



Cache Block size = 2^N bytes

Number of cache sets = 2^M

Offset = N bits




Set Index = M bits

Tag = 32 – (N + M) bits

Observation:

It is essentially unchanged from the direct-mapping formula


8. Set Associative Cache: Example



Offset, $N = 2$ bits

Block Number = $32 - 2 = 30$ bits

Check: Number of Blocks = 2^{30}



Number of Cache Blocks

= $4\text{KB} / 4\text{bytes} = 1024 = 2^{10}$

4-way associative, number of sets

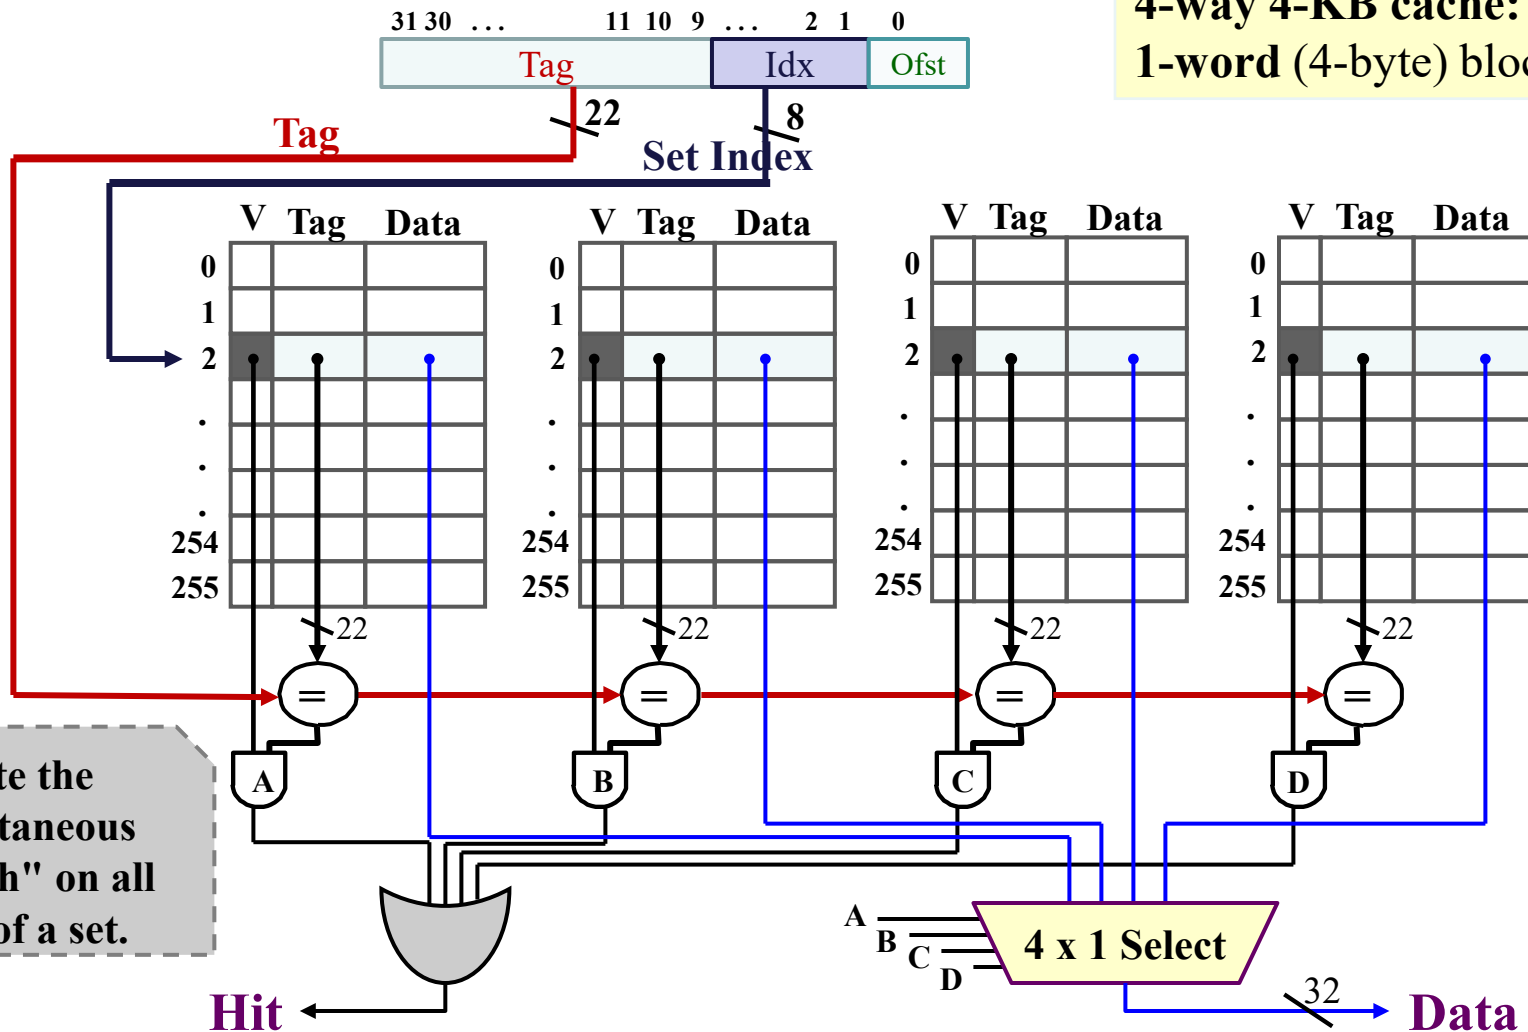
= $1024 / 4 = 256 = 2^8$

Set Index, $M = 8$ bits

Cache Tag = $32 - 8 - 2 = 22$ bits

8. Set Associative Cache: Circuitry

4-way 4-KB cache:
1-word (4-byte) blocks



8. SA Cache Example: Setup

■ Given:

- Memory access sequence: 4, 0, 8, 36, 0
- 2-way set-associative cache with a total of four 8-byte blocks → **total of 2 sets**
- Indicate hit/miss for each access



Offset, **N** = 3 bits

Block Number = $32 - 3 = 29$ bits

2-way associative, number of sets = $2 = 2^1$

Set Index, **M** = 1 bits

Cache Tag = $32 - 3 - 1 = 28$ bits

8. SA Cache Example: Load #1

Miss

4, 0, 8, 36, 0

Tag

Index Offset

■ Load from 4 →

00000000000000000000000000000000 0 100

Check: Both blocks in **Set 0** are invalid [**Cold Miss**]

Result: Load from memory and place in **Set 0 - Block 0**

Set Index	Block 0				Block 1			
	Valid	Tag	W0	W1	Valid	Tag	W0	W1
0	0 1	0	M[0]	M[4]	0			
1	0				0			

8. SA Cache Example: Load #2

Miss Hit
 4, **0**, 8, 36, 0

Tag

Index Offset

■ Load from 0 →

00000000000000000000000000000000 0 000

Result:

[Valid and Tags match] in **Set 0-Block 0** [**Spatial Locality**]

Set Index	Block 0				Block 1			
	Valid	Tag	W0	W1	Valid	Tag	W0	W1
0	1	0	M[0]	M[4]	0			
1	0				0			

8. SA Cache Example: Load #3

Miss	Hit	Miss
4, 0	,	8, 36, 0

Tag

Index Offset

■ Load from 8 →

00000000000000000000000000000000	1	000
----------------------------------	---	-----

Check: Both blocks in **Set 1** are invalid [**Cold Miss**]

Result: Load from memory and place in **Set 1 - Block 0**

Set Index	Block 0				Block 1			
	Valid	Tag	W0	W1	Valid	Tag	W0	W1
0	1	0	M[0]	M[4]	0			
1	0 1	0	M[8]	M[12]	0			

8. SA Cache Example: Load #4

Miss Hit Miss Miss

4	0	8	36	0
---	---	---	----	---

Index Offset

- Load from 36 ➡

[illegible]

0 | 100

Check: [Valid but tag mismatch] **Set 0 - Block 0**
[Invalid] **Set 0 - Block1** [**Cold Miss**]

Result: Load from memory and place in **Set 0 - Block 1**

Set Index	Block 0				Block 1			
	Valid	Tag	W0	W1	Valid	Tag	W0	W1
0	1	0	M[0]	M[4]	0 1	2	M[32]	M[36]
1	1	0	M[8]	M[12]	0			

8. SA Cache Example: Load #5

Miss	Hit	Miss	Miss	Hit
4	0	8	36	0
Tag			Index	Offset

■ Load from 0 →

00000000000000000000000000000000	0	000
----------------------------------	---	-----

Check: [Valid and tags match] **Set 0-Block 0**

[Valid but tags mismatch] **Set 0-Block 1**


[Temporal Locality]

Set Idx	Block 0				Block 1			
	Valid	Tag	W0	W1	Valid	Tag	W0	W1
0	1	0	M[0]	M[4]	1	2	M[32]	M[36]
1	1	0	M[8]	M[12]	0			

9. Cache Performance

Observations:

1. Cold/compulsory miss remains the same irrespective of cache size/associativity.
2. For the same cache size, conflict miss goes down with increasing associativity.
3. Conflict miss is 0 for FA caches.
4. For the same cache size, capacity miss remains the same irrespective of associativity.
5. Capacity miss decreases with increasing cache size.



Total Miss = Cold miss + Conflict miss + Capacity miss


Capacity miss (FA) = Total miss (FA) – Cold miss (FA), when Conflict Miss $\rightarrow 0$

10. Block Replacement Policy (1/3)

- **Set Associative or Fully Associative Cache:**
 - Can choose where to place a memory block
 - Potentially replacing another cache block if full
 - Need **block replacement policy**
- **Least Recently Used (LRU)**
 - **How:** For cache hit, record the cache block that was accessed
 - When replacing a block, choose one which has not been accessed for the longest time
 - **Why:** Temporal locality

10. Block Replacement Policy (2/3)

- Least Recently Used policy in action:
 - 4-way SA cache
 - Memory accesses: **0 4 8 12 4 16 12 0 4**



10. Block Replacement Policy (3/3)

- Drawback for LRU
 - Hard to keep track if there are many choices
- Other replacement policies:
 - First in first out (FIFO)
 - Random replacement (RR)
 - Least frequently used (LFU)

11. Summary: Cache Organizations

One-way set associative (direct mapped)

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

11. Summary: Cache Framework (1/2)

Block Placement: Where can a block be placed in cache?

Direct Mapped:

- Only one block defined by index

N-way Set-Associative:

- Any one of the **N** blocks within the set defined by index

Block Identification: How is a block found if it is in the cache?

Direct Mapped:

- Tag match with only one block

N-way Set Associative:

- Tag match for all the blocks within the set

11. Summary: Cache Framework (2/2)

Block Replacement: Which block should be replaced on a cache miss?

Direct Mapped:

- No Choice

N-way Set-Associative:

- Based on replacement policy

Write Strategy: What happens on a write?

Write Policy: Write-through vs write-back

Write Miss Policy: Write allocate vs write no allocate

12. Exploration: Improving Cache Penalty


Average Access Time

$$= \text{Hit rate} \times \text{Hit Time} + (1 - \text{Hit rate}) \times \text{Miss penalty}$$


- So far, we tried to improve Miss Rate:
 - Larger block size
 - Larger Cache
 - Higher Associativity
- What about **Miss Penalty**?

12. Exploration: Multilevel Cache

- Options:
 - Separate data and instruction caches, or a unified cache
- Sample sizes:
 - **L1**: 32KB, 32-byte block, 4-way set associative
 - **L2**: 256KB, 128-byte block, 8-way associative
 - **L3**: 4MB, 256-byte block, Direct mapped



12. Exploration: Intel Processors




Pentium 4 Extreme Edition

L1: 12KB I\$ + 8KB D\$

L2: 256KB

L3: 2MB



Itanium 2 McKinley


L1: 16KB I\$ + 16KB D\$

L2: 256KB

L3: 1.5MB – 9MB

12. Exploration: Trend: Intel Core i7-3960K

Intel® Core™ i7-3960X Processor Die Detail



Intel Core i7-3960K

per die:

-2.27 billion transistors

-15MB shared Inst/Data Cache (LLC)

per Core:

-32KB L1 Inst Cache

-32KB L1 Data Cache

-256KB L2 Inst/Data Cache

-up to 2.5MB LLC