


IT5002

# Computer Systems and Applications

## Lecture 12

### Process Management

[colintan@nus.edu.sg](mailto:colintan@nus.edu.sg)



# Introduction

- **In this lecture we will look at:**
  - The difference between a program and a process.
  - Interrupts and How They Work
  - Process States.
  - How to run Multiple Processes in a Single CPU
  - Process Creation, termination and zombies.

# Program vs. Process

- **A program consists of:**
  - Machine instructions (and possibly source code).
  - Data
  - A program exists as a file on the disk. E.g. command.exe, winword.exe
- **A process consists of:**
  - Machine instructions (and possibly source code).
  - Data.
  - Context
  - Exists as instructions and data in memory.
  - MAY be executing on the CPU.

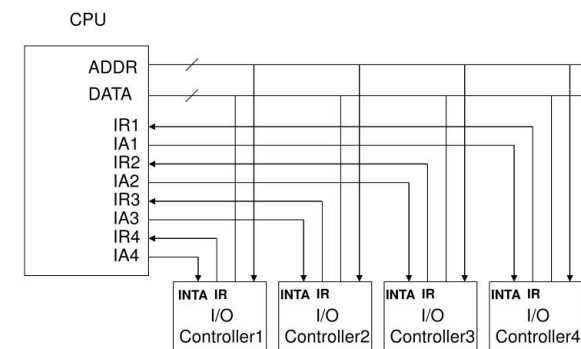
# Program vs. Process

- **A single program can produce multiple processes.**
  - E.g. chrome.exe is a single program.
  - But every tab in Chrome is a new process!

# Interrupts

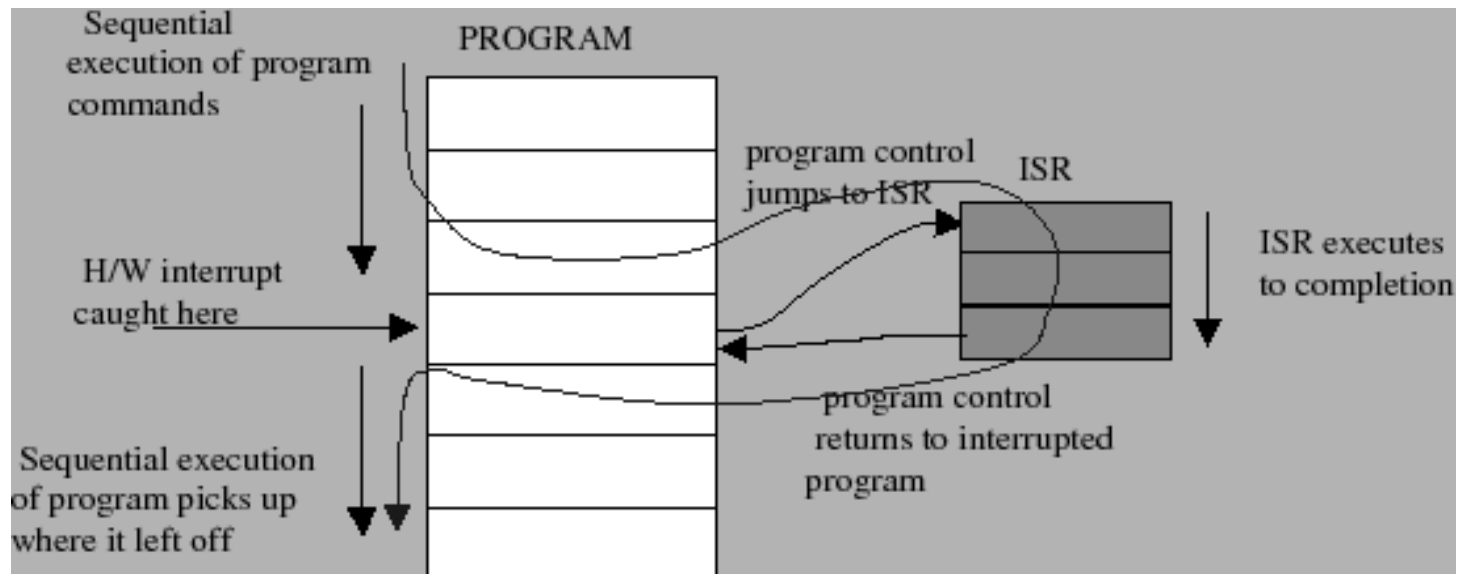
- **When a device needs attention from the CPU, it triggers what is called an “interrupt”:**
  - Each device is connected to the CPU via an input called an “Interrupt Request” or IRQ line.
  - ✓ When a device needs attention, it pull the IRQ line HIGH or LOW (dependin on whether line is “active high” or “active low”)

Multiple Interrupt Lines



# Interrupts

- This line is checked at the end of the WB stage in the CPU Execution Cycle.
- Diagrams below show the CPU's program execution flow when an interrupt occurs:



# Interrupts

- If line has been pulled, CPU interrupts the code it is currently running to run code to attend to the device:
  - ✓ **Current PC is pushed onto the stack.**
  - ✓ **CPU consults an “interrupt vector table” to look for the address of the “interrupt service routine” or ISR – a small bit of code that will read/write/tend to the device.**
  - ✓ **This address is loaded into PC, and the CPU starts executing the handler.**
  - ✓ **When the handler exits, the previous PC value is popped off the stack and back into PC, and execution resumes at the interrupted point.**

# Interrupts

- **The CPU asserts the interrupt acknowledge (IA) line to tell the device that its request has been handled.**
  - Sometimes the CPU will de-assert the IRQ line instead of employing a separate IA line.
- **Interrupts are key to allowing us to run multiple processes on a CPU:**
  - A hardware device called a “timer” will interrupt the CPU every millisecond.
  - Interrupt handler will switch to a new process.




# Execution Modes



- **Programs usually run sequentially.**
  - Each instruction is executed one after the other.
- **Having multiple cores or CPUs allow parallel (“concurrent”) execution.**
  - Streams of instructions with no dependencies are allowed to execute together.
- **A multitasking OS allows several programs to run “concurrently”.**
  - Interleaving, or “time-slicing”.

# Execution Modes

- 1 CPU: timesliced execution of tasks



- multiprocessor: timeslicing on n CPUs



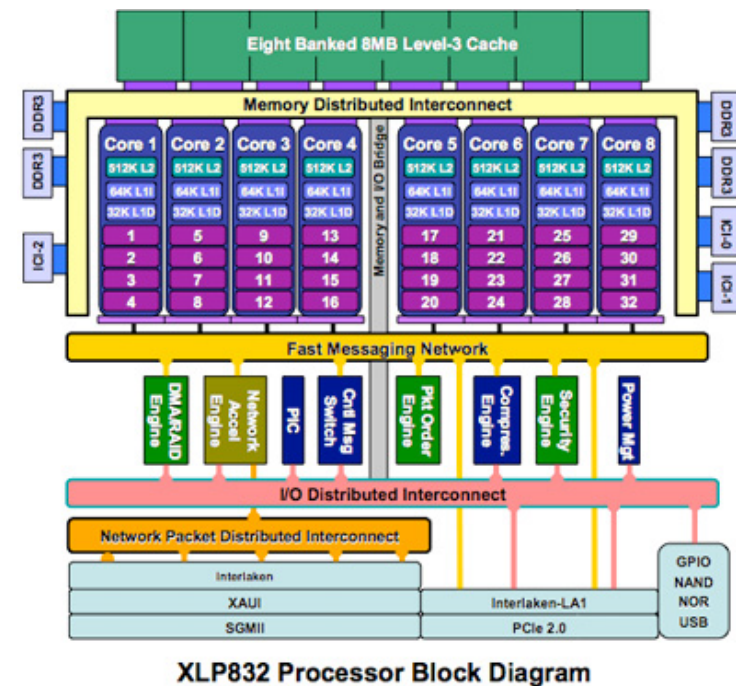
Note: we mostly assume no. processes  $\geq$  no. of CPU otherwise can have idle task

## Task Management

# PROCESSES AND PROCESS MANAGEMENT

# The Process Model


- In this lecture we will assume a single processor with a single core.
  - This is a legitimate assumption because in general the number of processes  $\gg$  the number of cores.
  - So each core must still switch between processes.



# The Process Model


- **Since we have only a single-core single processor:**
  - At any one time, at most one process can execute.
  - Figures (a) to (c) below illustrate what happens:

One program counter




(a)

Four program counters



(b)



(c)

# The Process Model

- **Figure (b) shows what “appears” to be happening in a single processor system running multiple processes:**
  - There are 4 processes each with its own program counter (PC) and registers.
  - All 4 processes run independently of each other at the same time.

# The Process Model

- **Figure (a) shows what actually happens.**
  - There is only a single PC and a single set of registers.
  - When one process ends, there is a “context switch” or “process switch”:
    - ✓ PC, all registers and other process data for Process A is copied to memory.
    - ✓ PC, register and process data for Process B is loaded and B starts executing, etc.
  - Figure (c) illustrates how processes A to D share CPU time.


# Process States

- **A process can be in one of 3 possible states:**
  - **Running**
    - ✓ The process is actually being executed on the CPU.
  - **Ready**
    - ✓ The process is ready to run but not currently running.
    - ✓ A “scheduling algorithm” is used to pick the next process for running.
  - **Blocked.**
    - ✓ The process is waiting for “something” to happen so it is not ready to run yet.
    - ✓ E.g. include waiting for inputs from another process.



# Process States


- The diagram below shows the 3 possible states and the transitions between them.



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

# Process States

- The figure below shows how the processes are organized.
  - The lowest layer selects (schedules) which process to run next.
    - ✓ This is subject to “scheduling policies” which we will look at in a later lecture.



# Switching between Processes

- **When a process runs, the CPU needs to maintain a lot of information about it. This is called the “process context”.**
  - CPU register values.
  - Stack pointers.
  - CPU Status Word Register.
    - ✓ **This maintains information about whether the previous instruction resulted in an overflow or a “zero”, whether interrupts are enabled, etc.**
    - ✓ **This is needed for branch instructions – assembly equivalents of “if” statements.**

The AVR Status Register – SREG – is defined as:

Bit	7	6	5	4	3	2	1	0
0x3F (0x5F)	I	T	H	S	V	N	Z	C
Read/Write	R/W	R/W	R/W	R/W	R/W	R/W	R/W	R/W
Initial Value	0	0	0	0	0	0	0	0

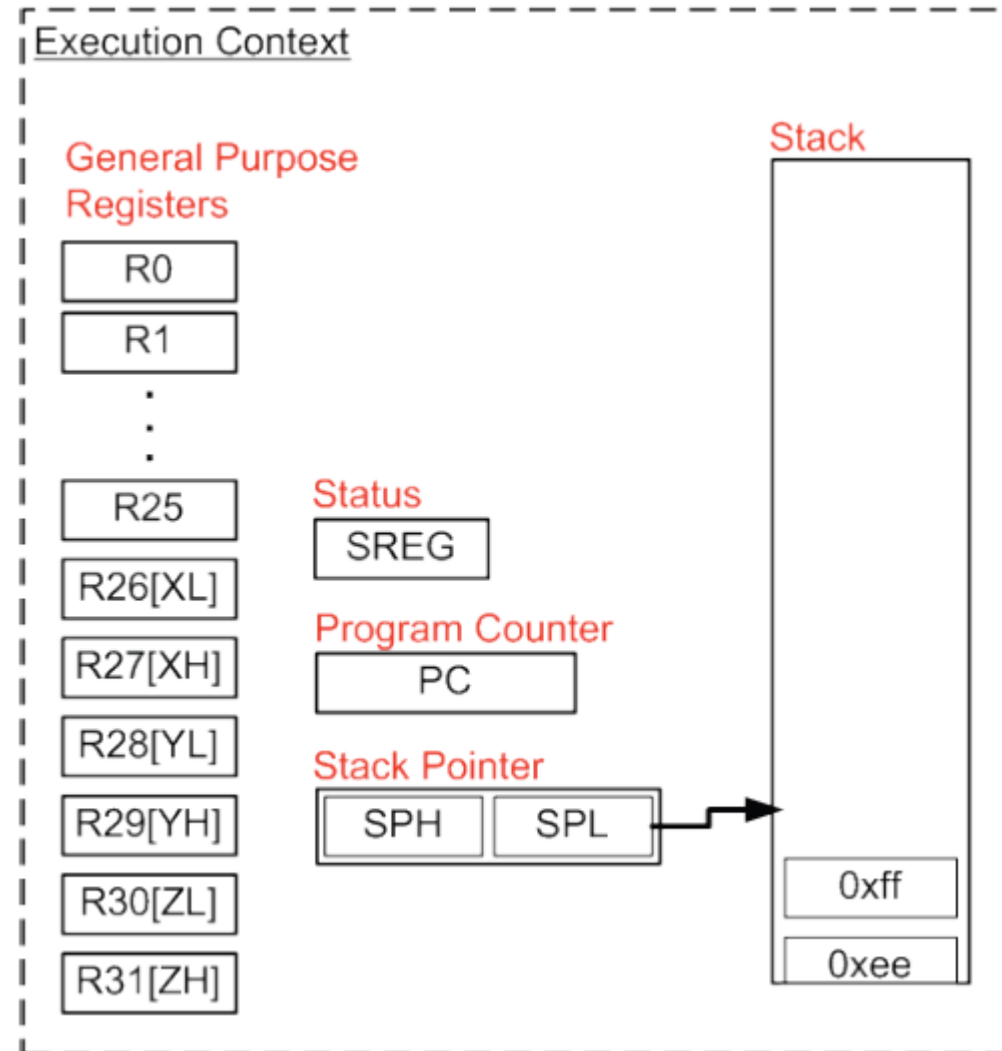
SREG

# Switching between Processes

- All of these values change as a process runs.
- When a process is blocked or put into a **READY** state, a new process will be picked to take control of the CPU.
  - All the information for the current process must be saved!
  - The information for the new process must be loaded into the registers, stack pointer and status registers!
    - ✓ This is to allow the new process to run like as though it was never interrupted!
- This process is known as “context switching”.

# Context Switching on the FreeRTOS Atmega Port

- **Each process is allocated a stack.**
  - Exactly what you learnt in IT5003
  - We use the term “task” instead of process here but means the same thing.
- **The diagram shows the complete Atmega context.**
  - Registers R0-R31, PC.
  - Status register SREG.
  - Stack pointer SPH/SPL.



# Context Switching on the FreeRTOS Atmega Port

- FreeRTOS implements context saving in a macro called “portSAVE\_CONTEXT”.**

```
#define portSAVE_CONTEXT() \
asm volatile ( \
    "push r0 \n\t" \
    "in r0, __SREG__ \n\t" \
    "cli \n\t" \
    "push r0 \n\t" \
    "push r1 \n\t" \
    "clr r1 \n\t" \
    "push r2 \n\t" \
    ... \
    "push r31 \n\t" \
    "in r26, __SP_L__ \n\t" \
    "in r27, __SP_H__ \n\t" \
    "sts pxCurrentTCB+1, r27 \n\t" \
    "sts pxCurrentTCB, r26 \n\t" \
    "sei \n\t" : : : \
);
```

// Save R0

// Read in status register SREG to R0

// Disable all interrupts for atomicity

// Save SREG

// Save R1

// AVR C expects R1 to be 0, so clear it.

// Save R2 to R31

// Read in stack pointer low byte

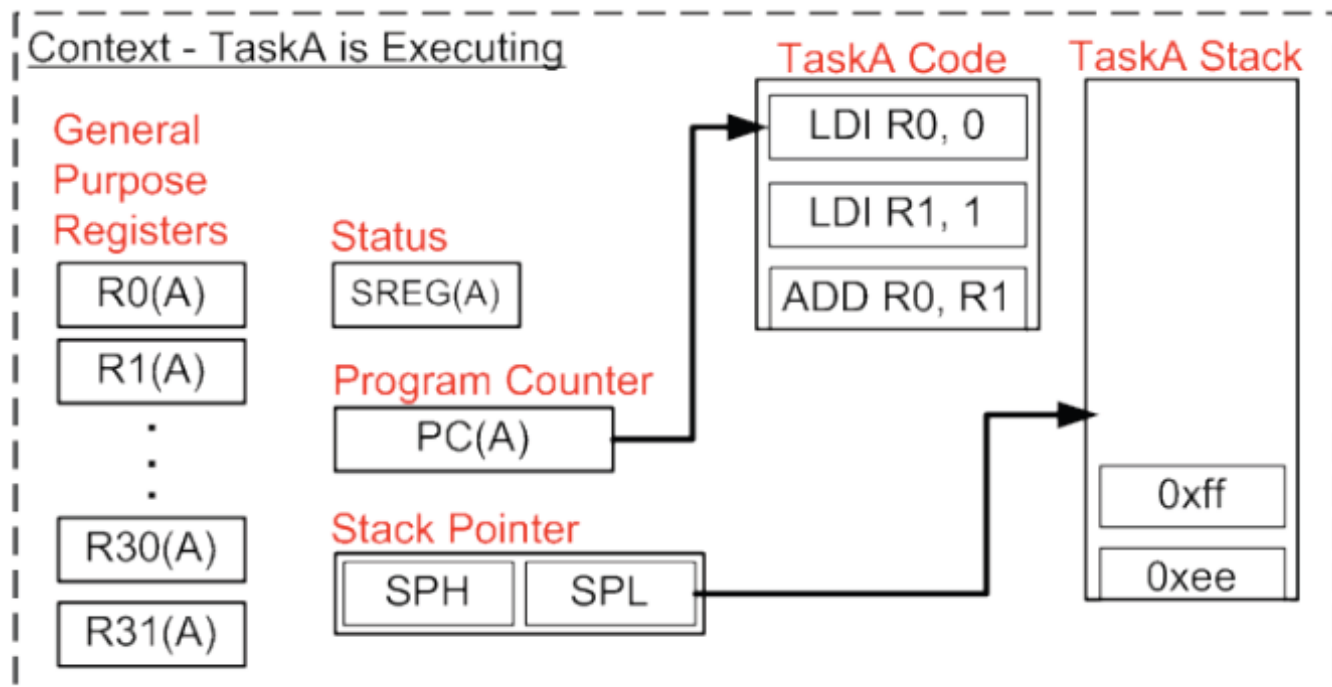
// and high byte

// And save it to pxCurrentTCB

// Re-enable interrupts

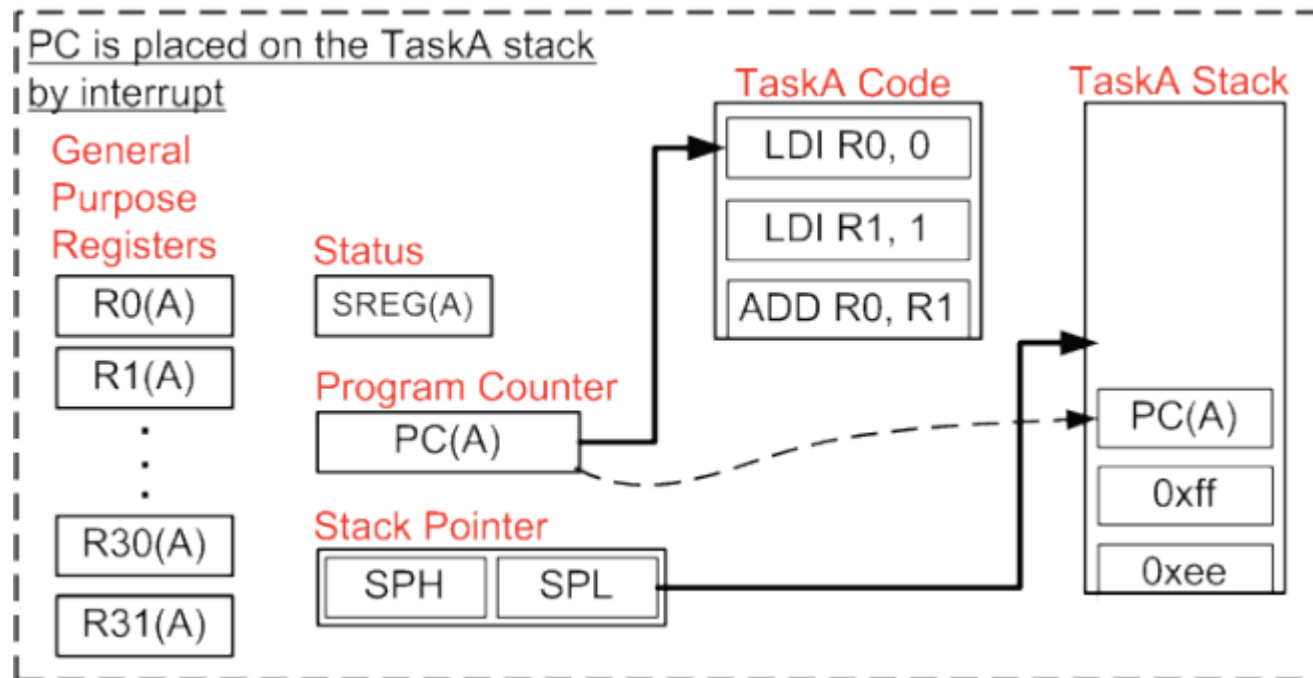
# Context Switching on the FreeRTOS Atmega Port

- We will now see step-by-step how this works.
- Assume that at first Task A is executing.
  - PC would be pointing at Task A code, SPH/SPL pointing at Task A stack, Registers R0-R31 contain Task A data.



# Context Switching on the FreeRTOS Atmega Port

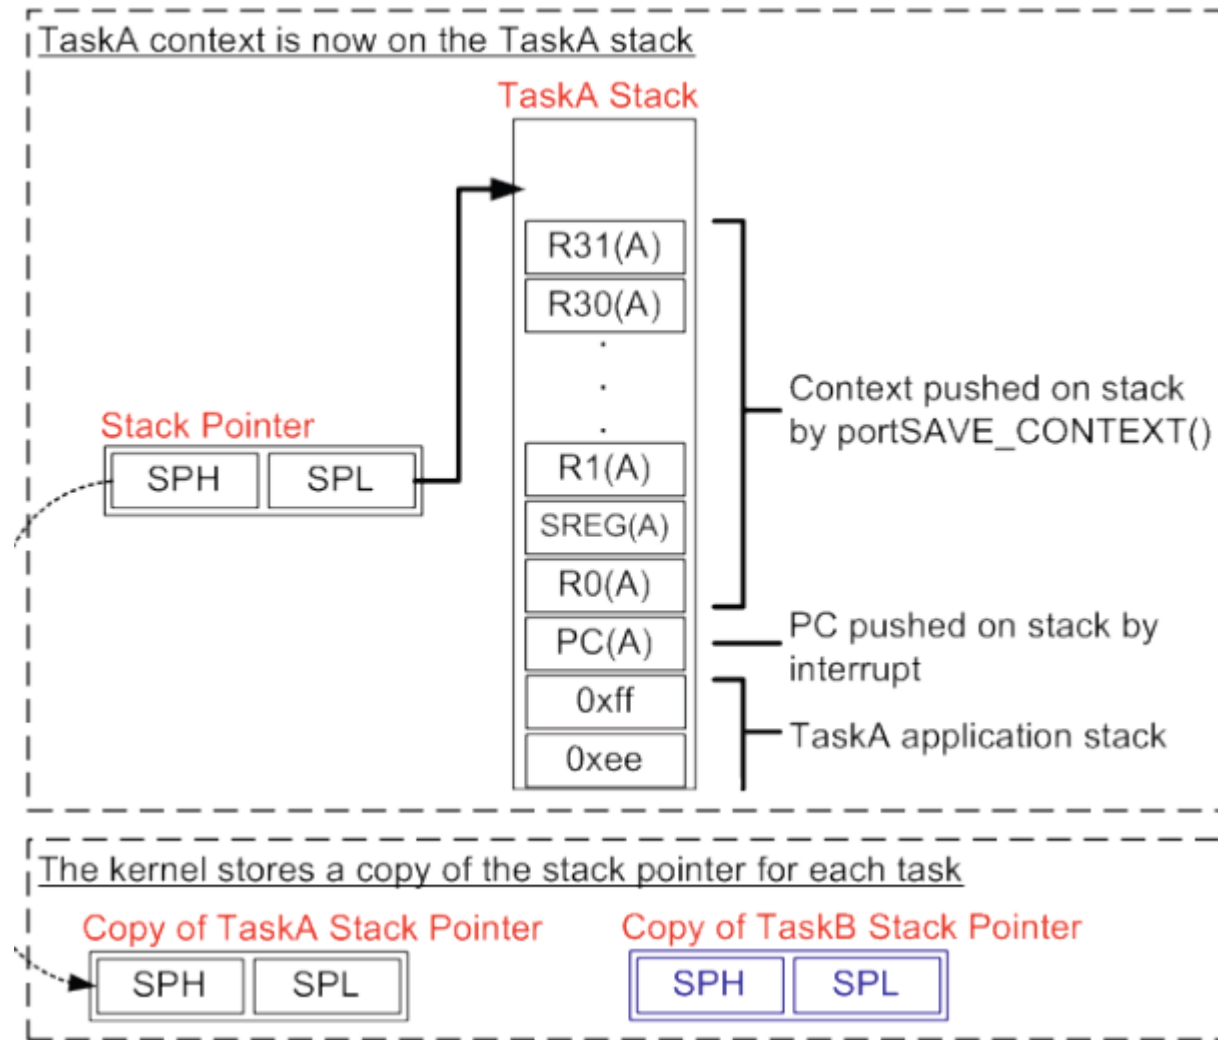
- FreeRTOS relies on regular interrupts from Timer 0 every millisecond to switch between tasks. When the interrupt triggers, PC is placed onto Task A's stack.**





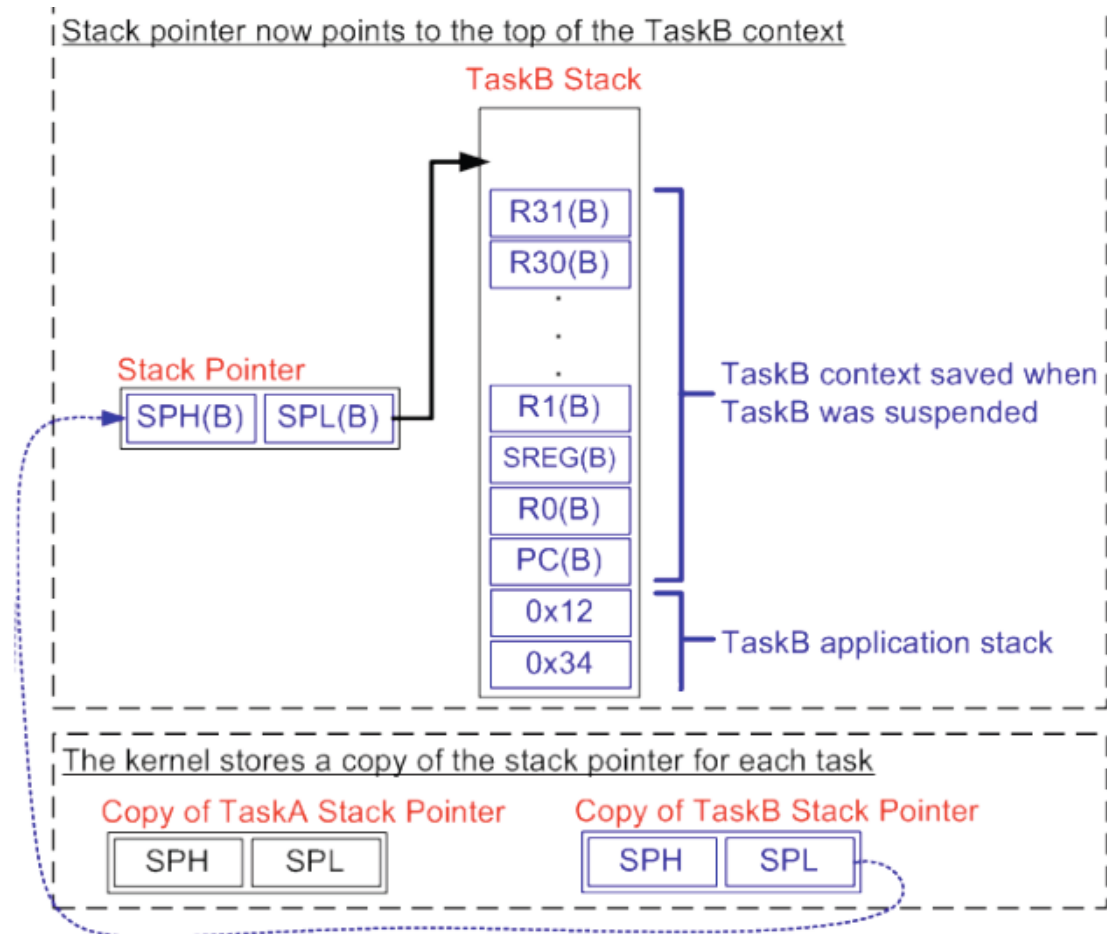
# Context Switching on the FreeRTOS Atmega Port

- The ISR calls `portSAVECONTEXT`, resulting in Task A's context being pushed onto the stack.
- `pxCurrentTCB` will also hold SPH/SPL after the context save.
  - This must be saved by the kernel.



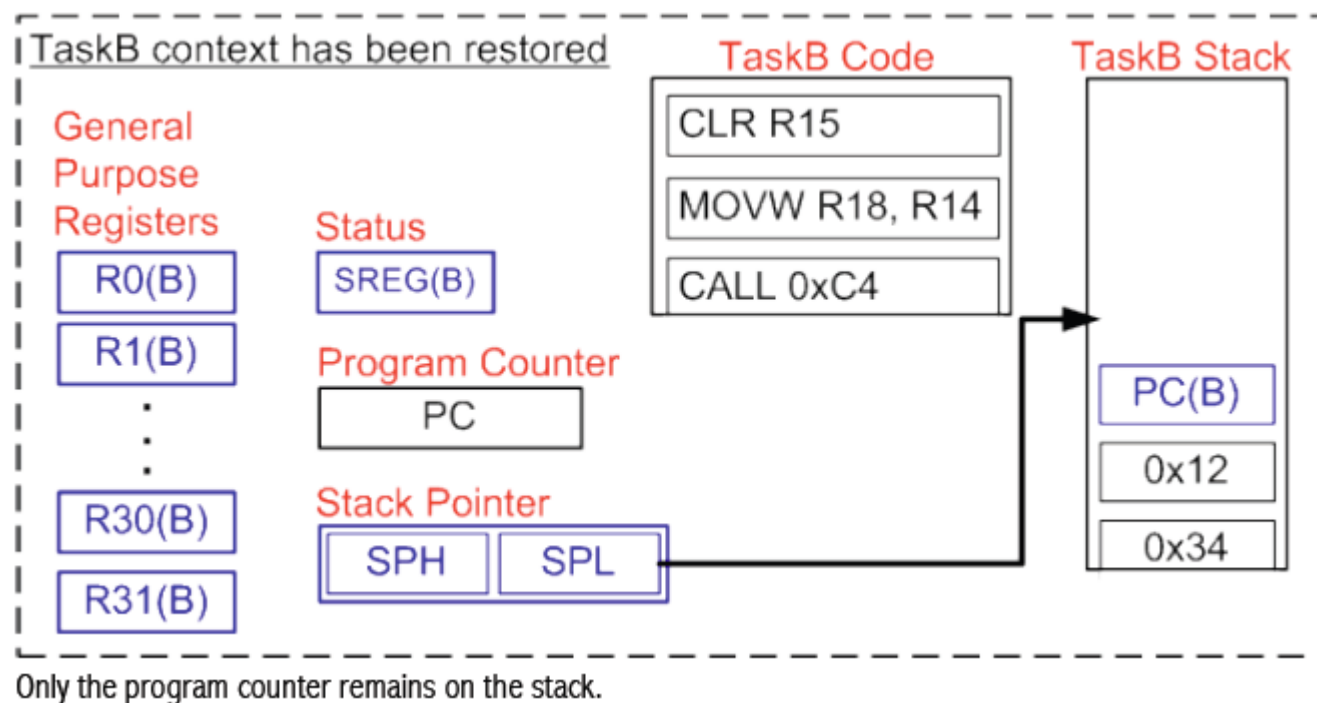
# Context Switching on the FreeRTOS Atmega Port

- **The kernel then selects Task B to run, and copies its SPH/SPL values into pxCurrentTCB and calls portRESTORE\_CONTEXT.**
  - The first two lines will copy pxCurrentTCB into SPH/SPL, causing SP to point to Task B's stack.



# Context Switching on the FreeRTOS Atmega Port

- **The rest of portRESTORE\_CONTEXT is executed, causing Task B's data to be loaded into R31-R0 and SREG.**
  - Now Task B can resume like as though nothing happened!



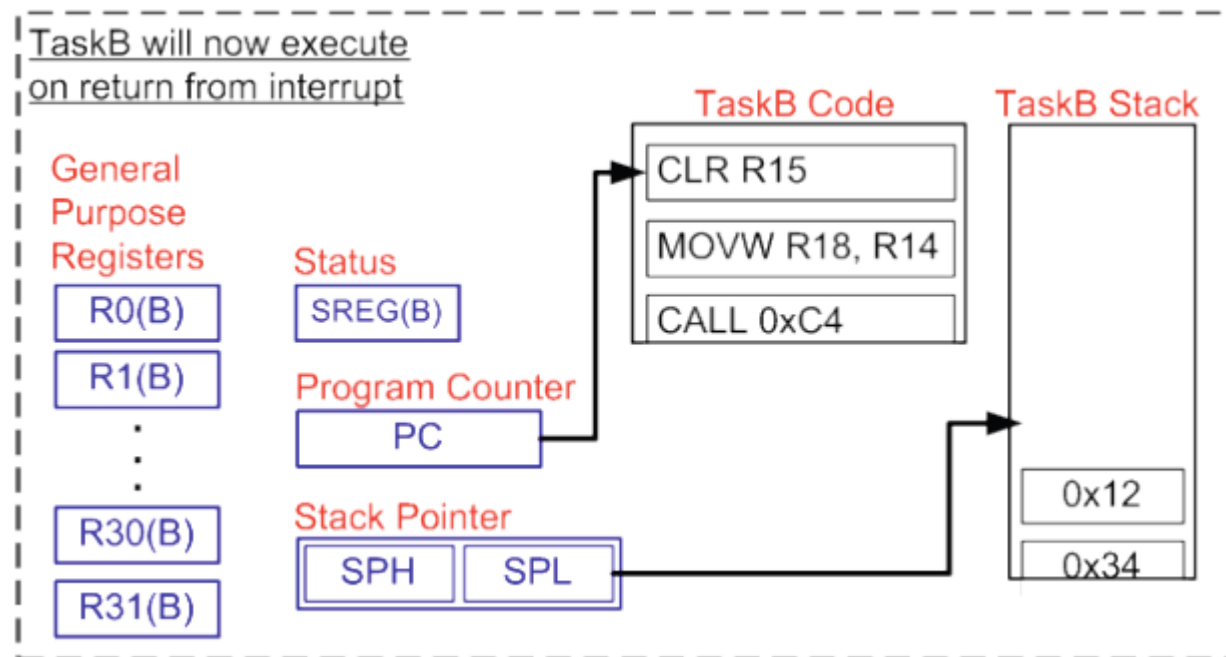
# Context Switching on the FreeRTOS Atmega Port

- **The reverse operation is portRESTORE\_CONTEXT. The stack pointer for the process being restored must be in pxCurrentTCB.**

```
#define portRESTORE_CONTEXT()\nasm volatile (\n    "out __SP_L__, %A0 \\n\\t"\\    // Copy SP_L and SP_H from the\n    "out __SP_H__, %B0 \\n\\t"\\    // pxCurrentTCB variable.\n    "pop r31 \\n\\t"\\            // Restore registers r31 to r1.\n    ...\n    "pop r0 \\n\\t"\\                // Pop out SREG\n    "out __SREG__, r0\\n\\t"\\        // And restore it.\n    "pop r0 \\n\\t": : "r" (pxCurrentTCB):\\    // Restore R0\n\n);
```

# Context Switching on the FreeRTOS Atmega Port

- **Only Task B's PC remains on the stack. Now the ISR exits, causing this value to be popped off onto the AVR's PC.**
  - PC points to the next instruction to be executed.
  - End result: Task B resumes execution, with all its data and SREG intact!



# Context Switching on the FreeRTOS Atmega Port

- **Here we looked at context switching controlled by a timer.**
  - **It can also be triggered by other things:**
    - ✓ **Currently running processed waiting for input.**
    - ✓ **Currently running task blocking on a synchronization mechanism (see next lecture).**
    - ✓ **Currently running task wants to sleep for a fixed period.**
    - ✓ **Higher priority task becoming “READY”.**
    - ✓ **...**

# Process Creation

- A process can be created in Python by using a `fork()` call:

```
# Python code to create child process
import os

def parent_child():
    n = os.fork()

    # n greater than 0 means parent process
    if n > 0:
        print("Parent process and id is : ", os.getpid())

    # n equals to 0 means child process
    else:
        print("Child process and id is : ", os.getpid())

# Driver code
parent_child()
```

- The creating process is called a “parent” process, while the created process is called a “child” process.

# Process Creation

- **When you run a program in your OS shell, the shell uses the OS to create a new process, then run the program in the new process.**
- **In Python this is done by using `subprocess.call`:**

```
import subprocess  
  
subprocess.call(["ls", "-lha"])
```

  - The launched is thus a child process of the shell.
- **Ultimately, all UNIX processes are children of “init”, the main starting process in UNIX.**

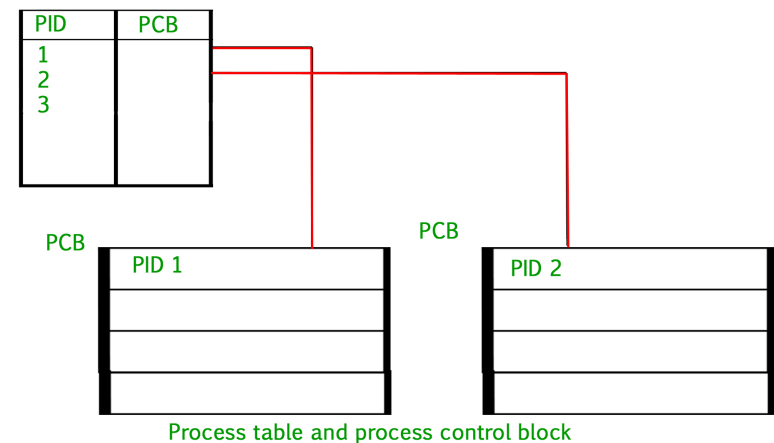


# Process Control Blocks

- When a process is created, the Operating System also creates a data structure to maintain information about that process:

- Called a “Process Control Block” (PCB) and contains:

- ✓ Process ID (PID)
- ✓ Stack Pointer
- ✓ Open files
- ✓ Pending signals
- ✓ CPU usage
- ✓ ...



- PCB is stored in a table called a “Process Table”.

- ✓ One Process Table for entire system.
- ✓ One PCB per process.

# Terminating A Process

- **When a process terminates:**
  - Most resources like open files, etc., can be released and returned to the system.
  - However the PCB is retained in memory:
    - ✓ **Allows child processes to return results to the parent.**
  - Parent retrieves the results using a “wait” function call, afterwhich the PCB is released.
- **What if the parent never calls “wait”?**
  - PCB remains in memory indefinitely.
  - Child becomes a “zombie”. Eventually process table will run out of space and no new process can be created.