


# IT5002

## Computer Systems and Applications

### Lecture 13

### Process Scheduling



# What does your Computer Spend its Time Doing?

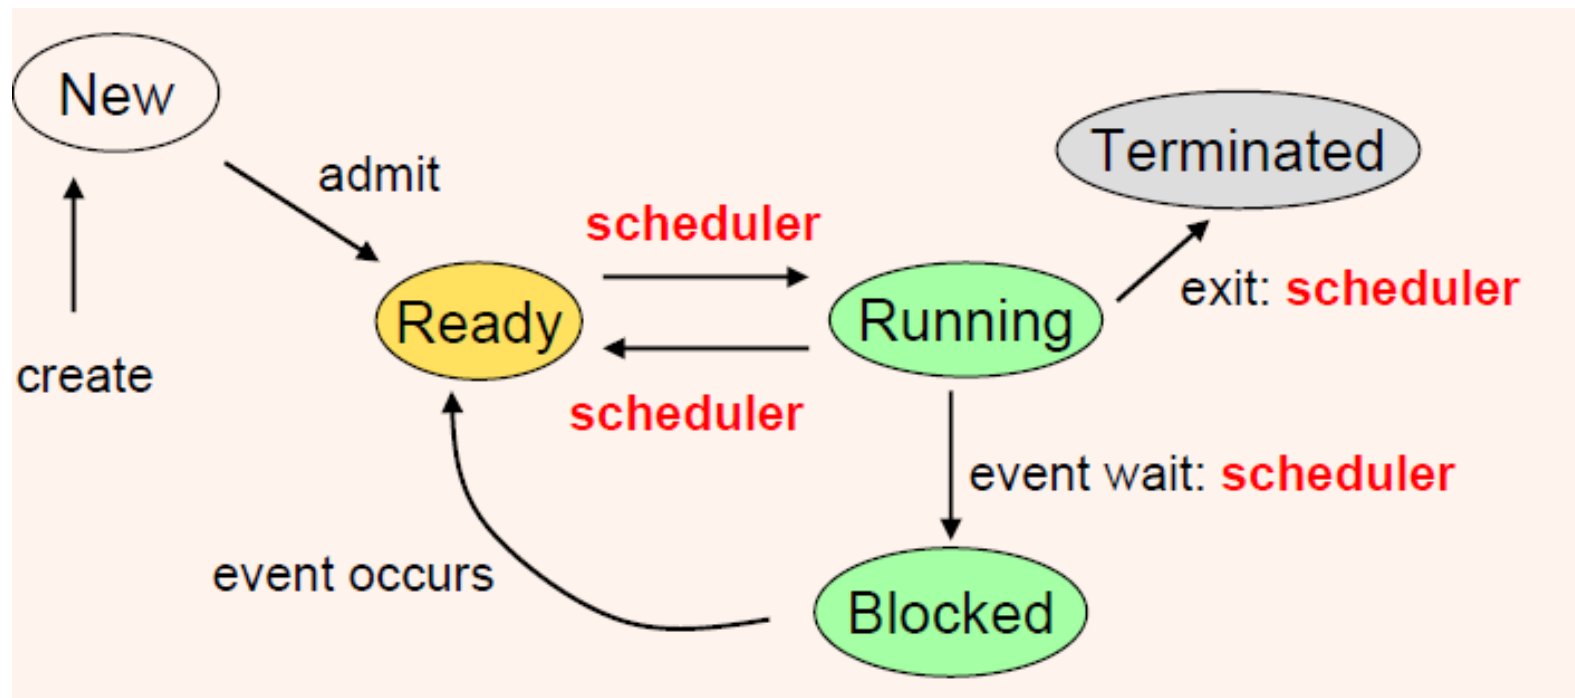
- **A mix of jobs:**
  - Computations + reading/writing memory.
  - Input/Output
    - ✓ Reading from the keyboard
    - ✓ Writing to the screen
    - ✓ Reading from the mouse
    - ✓ Sending/receiving data over the network.
    - ✓ Reading/writing the disk
    - ✓ ...
  - How do we manage all these varied jobs?

# The Scheduling Environment

- **Processes can be:**
  - **CPU bound**
    - ✓ Most of the time spent on processing on CPU
    - ✓ Graphics-intensive applications are considered to be “CPU” bound.
    - ✓ Multitasking opportunities come from having to wait for processing results.
  - **I/O bound**
    - ✓ Most of the time is spent on communicating with I/O devices
    - ✓ Multitasking opportunities come from having to wait for data from I/O devices.

# Process States

- Processes switch between a fixed set of states depending on events that take place.
  - Scheduler is invoked at various points as shown below.



# Generic Scheduler Algorithm

```
schedule() {  
    while (queue not empty) {  
        task = pick task from ready queue; // policy dependent  
        delete task from queue;  
        switch to task; // how is it this done? architecture dependent  
    }  
}
```

Question: How do we determine policies to pick the next task?

# Types of Multitaskers

- **Policies are determined by the kind of multitasking environment.**
  - **Batch Processing**
    - ✓ **Not actually multitasking since only one process runs at a time to completion.**
  - **Co-operative Multitasking**
    - ✓ **Currently running processes cannot be suspended by the scheduler.**
    - ✓ **Processes must volunteer to give up CPU time.**
  - **Pre-emptive Multitasking**
    - ✓ **Currently running processes can be forcefully suspended by the scheduler.**

# Types of Multitaskers

- Real-Time Multitasking

- ✓ Processes have fixed deadlines that must be met.

- What if we don't meet the deadlines?

- ✓ Hard Real Time Systems: Disaster strikes! System fails, possibly catastrophically!

- ✓ Soft Real Time Systems: Mostly just an inconvenience. Performance of system is degraded.



Photo Copyright Grain Murray

AIRLINERS.NET

# Scheduling Policies for Multitaskers

- **Scheduling Policies enforce a priority ordering over processes.**
  - As mentioned earlier, determined by multitasking type.
- **Example Policies**
  - Simplest Policy (Great for all types of multitaskers)
    - ✓ **Fixed Priority**
  - Policies for Batch Processing
    - ✓ **First-come First Served (FCFS)**
    - ✓ **Shortest Job First (SJF)**
  - Policies for Co-operative Multitaskers
    - ✓ **Round Robin with Voluntary Scheduling (VS)**



# Scheduling Policies for Multitaskers

- **Example Policies**
  - Policies for Pre-emptive Multitaskers
    - ✓ **Round Robin with Timer (RR)**
    - ✓ **Shortest Remaining Time (SRT)**
  - Policies for Real-Time Multitaskers (Not covered)
    - ✓ **Rate Monotonic Scheduling (RMS)**
    - ✓ **Earliest Deadline First Scheduling (EDF)**

# Fixed Priority Policy

- **This is a simple policy that can be used across any type of multitasker.**
  - Each task is assigned a priority by the programmer.
    - ✓ Usually priority number 0 has the highest priority.
  - Tasks are queued according to priority number.
  - Batch, Co-operative:
    - ✓ Task with highest priority is picked to be run next.
  - Pre-emptive, Real-Time:
    - ✓ When a higher priority task becomes ready, current task is suspended and higher priority task is run.

# Batch Scheduling Policies

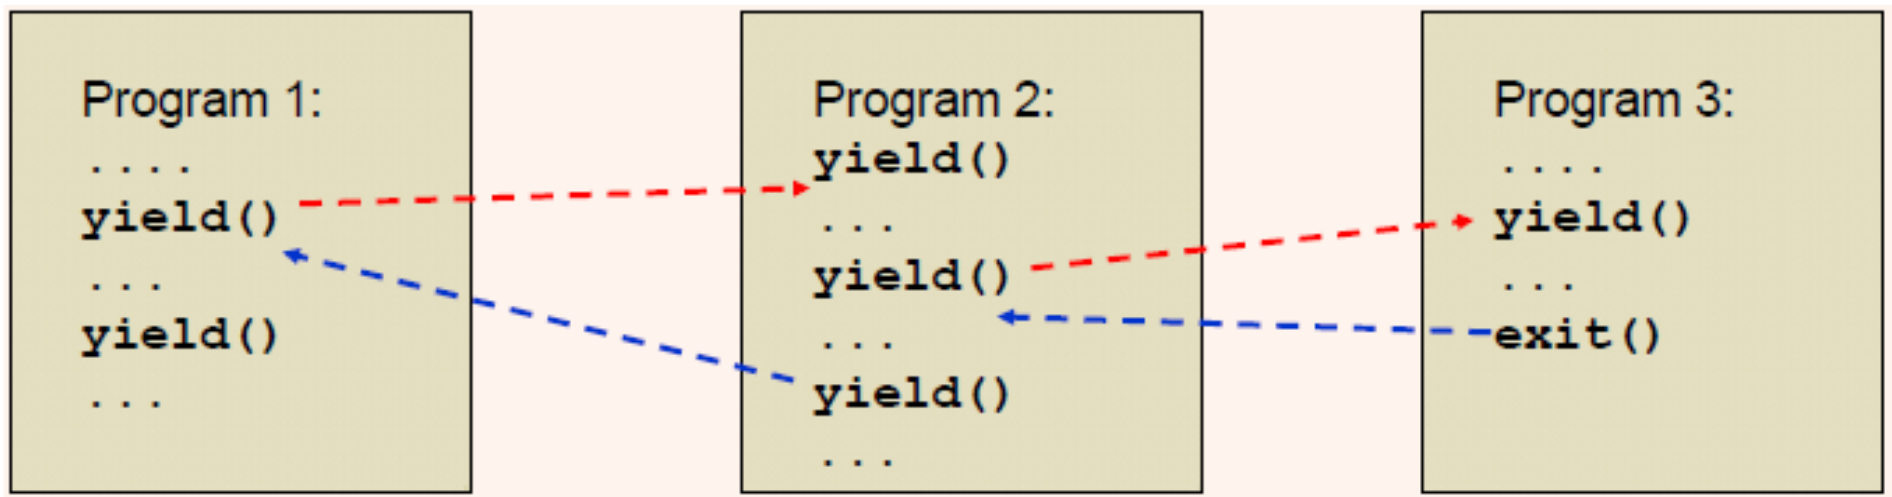
- **First Come First Served**
  - Arriving jobs are stored in a queue.
  - Jobs are removed in turn and run.
  - Particularly suited for batch systems.
  - Extension for interactive systems:
    - ✓ Jobs removed for running are put back into the back of the queue.
    - ✓ This is also known as “round-robin scheduling”
  - Starvation free as long as earlier jobs are bounded.

# Batch Scheduling Policies

- **Shortest Job First**
  - Processes are ordered by total CPU time used.
  - Jobs that run for less time will run first.
  - Reduces average waiting time if number of processes is fixed.
  - Potential for starvation.

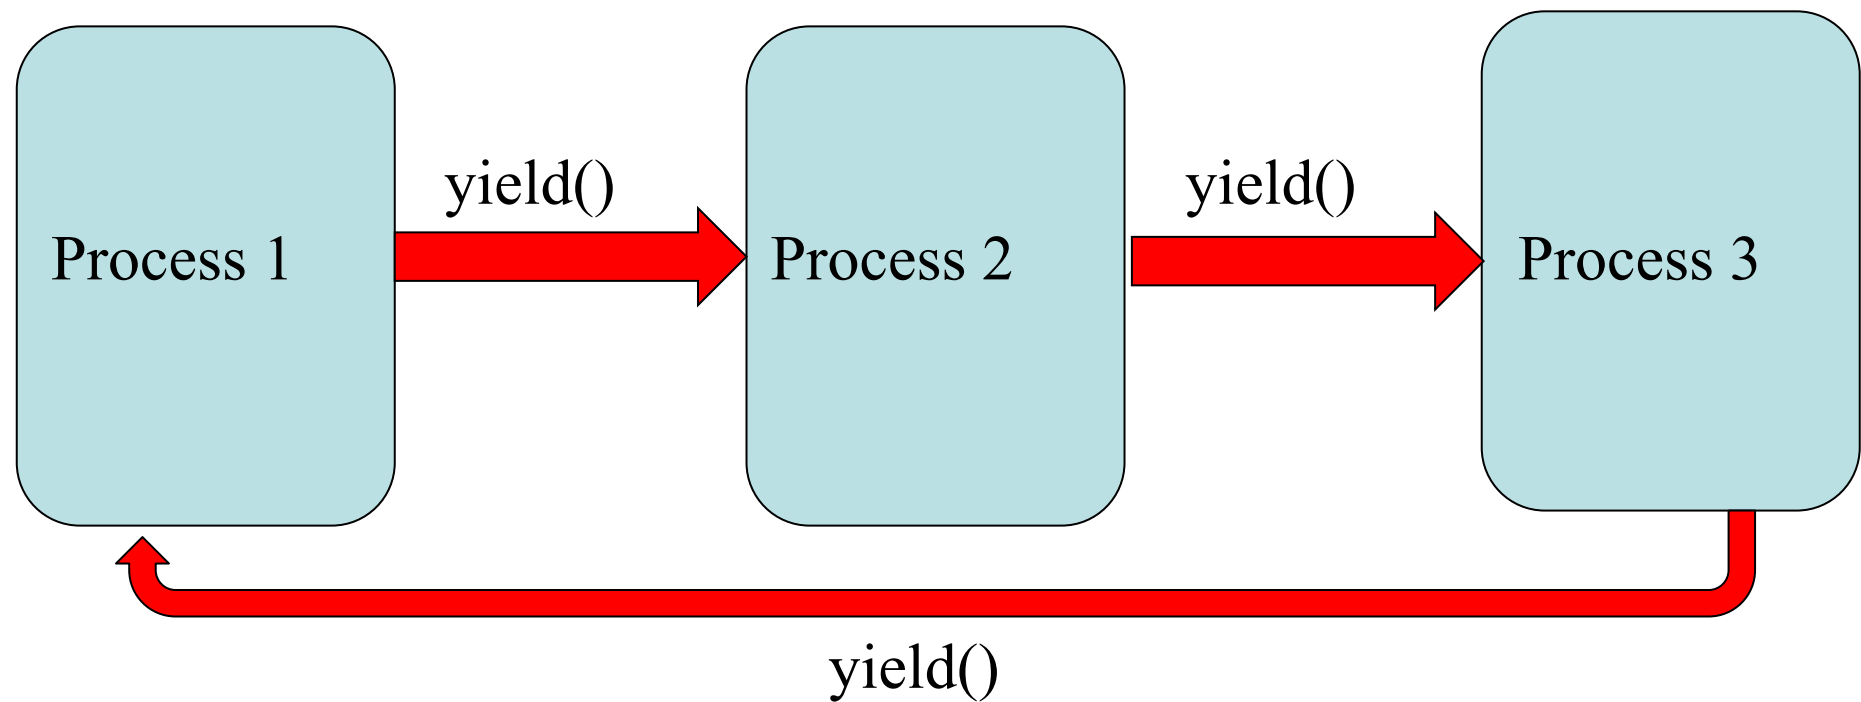
# Co-operative Scheduling Policies

- **Voluntary Scheduling.**
  - Processes call a special “yield” function.
    - ✓ **This invokes the scheduler.**
    - ✓ **Causes the process to be suspended and another process started up.**



# Co-operative Scheduling Policies

- In many systems VS is used with a round-robin arrangement.



# Pre-emptive Scheduling Policies

- **Shortest Remaining Time**
  - Pre-emptive form of SJF.
  - Processes are ordered according to remaining CPU time left.
- **Round-robin with Timer**
  - Each process is given a fixed time slot  $c_i$ .
  - After time  $c_i$ , scheduler is invoked and next task is selected on a round-robin basis.

# Managing Multiple Policies

- **Multiple policies can be implemented on the same machine using multiple queues:**
  - Each queue can have its own policy.
  - This scheme is used in Linux, as we will see shortly.

high priority:	P1, P3	(RR policy)
medium priority:	P2	(RR policy)
low priority:	P4, P5	(batch queue, FCFS policy)



# Scheduling in Linux

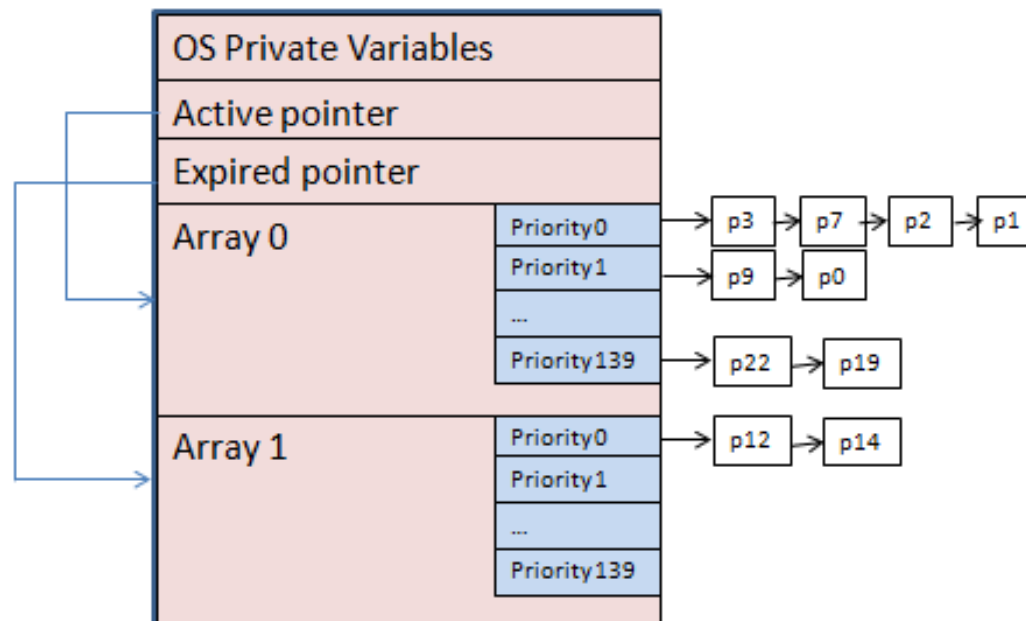
- **Processes in Linux are dynamic:**
  - New processes can be created with `fork()`
  - Existing processes can exit.
- **Priorities are also dynamic:**
  - Users and superusers can change priorities using “nice” values.
  - `nice -n 19 tar cvzf archive.tgz *`
    - ✓ Allows tar to run with a priority lowered by 19 to reduce CPU load.
    - ✓ Normal users can only  $0 \leq n \leq 19$
    - ✓ Superusers can specify  $-20 \leq n \leq 19$ . Negative nice increases priority.

# Scheduling in Linux

- **Linux maintains three types of processes:**
  - **Real-time FIFO:**
    - ✓ **RT-FIFO processes cannot be pre-empted except by a higher priority RT-FIFO process.**
  - **Real-time Round-Robin:**
    - ✓ **Like RT-FIFO but processes are pre-empted after a time slice.**
  - **Linux only has “soft real-time” scheduling.**
    - ✓ **Cannot guarantee deadlines, unlike RMS and EDF we saw earlier.**
    - ✓ **Priority levels 0 to 99**
  - **Non-real time processes**
    - ✓ **Priority levels 100 to 139**

# Scheduling in Linux

- **Linux maintains 280 queues in two sets of 140:**
  - An active set.
  - An expired set.



# Scheduling in Linux

- **The scheduler is called at a rate of 1000 Hz.**
  - E.g. time tick is 1 ms, called a “jiffy”.
  - RT-FIFO processes are always run if any are available.
  - Otherwise:
    - ✓ Scheduler picks highest priority process in active set to run.
    - ✓ When its “time quantum” is expired, it is moved to the expired set. Next highest priority process is picked.
    - ✓ When active set is empty, active and expired pointers are swapped. Active set becomes expired set and vice versa.
    - ✓ Scheme ensures no starvation of lowest priority processes.

# Scheduling in Linux

- **What happens if a process becomes blocked? (e.g. on I/O)**
  - CPU time used so far is recorded. Process is moved to a queue of blocked processes.
  - When process becomes runnable again, it continues running until its time quantum is expired.
  - It is then moved to the expired set.
- **When a process becomes blocked its priority is often upgraded (see later).**

# Scheduling in Linux

- **Time quantum for RR processes:**
  - **Varies by priority. For example:**
    - ✓ **Priority level 100 – 800 ms**
    - ✓ **Priority level 139 – 5 ms**
    - ✓ **System load.**
- **How process priorities are calculated:**
  - **Priority = base + f(nice)+g(cpu usage estimate)**
    - ✓ **f(.) = priority adjustment from nice value.**
    - ✓ **g(.) = Decay function. Processes that have already consumed a lot of CPU time are downgraded.**

# Scheduling in Linux

- Other heuristics are used:
  - ✓ Age of process.
  - ✓ More priority for processes waiting for I/O – I/O boost.
  - ✓ Bias towards foreground tasks.
- I/O Boost:
  - Rationale:
    - ✓ Tasks doing read() has been waiting for a long time. May need quick response when ready.
    - ✓ Blocked/waiting processes have not run much.
    - ✓ Applies also to interactive processes – blocked on keyboard/mouse input.

# Scheduling in Linux

- Implementation: We can -
  - ✓ Boost time quantum.
  - ✓ Boost priority.
  - ✓ Do both.
- How long does this boost last?
  - ✓ Temporary boost for sporadic I/O
  - ✓ Permanent boost for the chronically I/O bound?
  - ✓ E.g. Linux gives -5 boost for interactive processes.